

# hw05. syswrite

406410114 郭晏誠

1. 請將 sys\_write 追蹤完畢，說明從剛開始到「於 terminal 中印出字串」為止的函數。

(1)b sys\_write 後停在 635 行 在 638 行設置中斷點

The screenshot shows a debugger interface with several tabs at the top: Console, Registers, Problems, Executables, Debugger Console, Memory, OS Resources, and another tab that is partially visible. Below the tabs, there is a stack trace:

```
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
arg buf = 39666576
arg count = 1
[1] from 0xffffffff81d5d672 in system_call+114 at arch/x86/kernel/entry_64.S:361
(no arguments)
[+]
Threads
[1] id 1 from 0xffffffff81290912 in SyS_write+20 at fs/read_write.c:635
```

At the bottom, there is a message about a breakpoint:

```
Breakpoint 1, SyS_write (fd=1, buf=39666576, count=1) at fs/read_write.c:635
635     SYSCALL_DEFINE3(write, unsigned int, fd, const char __user *, buf,
>>>
```

(2)按 C 跳到 638 用 s 開始追蹤

The screenshot shows a debugger interface with several tabs at the top: main.c, process.c, read\_write.c, 0xffff0, Registers, Problems, Executables, Debugger Console, Memory, OS Resources, and Perf Profile View. On the right side, there is a variables table:

Name	Type
fd	un
buf	co
count	siz
f	sti
ret	ss

Below the tabs, there is a stack trace:

```
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff81290945 in SYSC_write+19 at fs/read_write.c:638
SYSC_write (fd=1, buf=0x2481390 "\nmnt/sharedFolder/hw03.syscall.asm # /\033[J", count=1) at fs/read_write.c:638
638     struct fd f = fdget_pos(fd);
>>> █
```

(3)追蹤過程見下方表格，詳細截圖過程請見附錄

Index	File and column	Function code	Mark
1	fs/read_write.c:638	<code>struct fd f = fdget_pos(fd);</code>	
2	fs/read_write.c:270	<code>return __to_fd(__fdget_pos(fd));</code>	
3	fs/file.c:708	<code>unsigned long v = __fdget(fd);</code>	
4	fs/file.c:697	<code>return __fget_light(fd, FMODE_PATH);</code>	
5	fs/file.c:680	<code>struct files_struct *files = current-&gt;files;</code>	
6	./arm/x86/include/asm/current.h:14	<code>return this_cpu_read_stable(current_task);</code>	
7	fs/file.c:680	<code>struct files_struct *files = current-&gt;files;</code>	
8	fs/file.c:683	<code>if (atomic_read(&amp;files-&gt;count) == 1) {</code>	
9	./arm/x86/include/asm/atomic.h:27	<code>return ACCESS_ONCE((v)-&gt;counter);</code>	
10	fs/file.c:684	<code>file = __fcheck_files(files, fd);</code>	
11	include/linux/fdtable:77	<code>struct fdtable *fdt = rcu_dereference_raw(files-&gt;fdt);</code>	
12	include/linux/fdtable:79	<code>if (fd &lt; fdt-&gt;max_fds)</code>	
13	include/linux/fdtable:80	<code>return rcu_dereference_raw(fdt-&gt;fd[fd]);</code>	
14	fs/file.c:685	<code>if (!file    unlikely(file-&gt;f_mode &amp; mask))</code>	
15	fs/file.c:687	<code>return (unsigned long)file;</code>	
16	fs/file.c:709	<code>struct file *file = (struct file *) (v &amp; ~3);</code>	
17	fs/file.c:711	<code>if (file &amp;&amp; (file-&gt;f_mode &amp; FMODE_ATOMIC_POS)) {</code>	
18	fs/file.c:717	<code>return v;</code>	
19	include/linux/file.h:50	<code>return (struct fd){(struct file *) (v &amp; ~3), v &amp; 3};</code>	
20	fs/read_write.c:639	<code>ssize_t ret = -EBADF;</code>	
21	fs/read_write.c:641	<code>if (f.file) {</code>	
22	fs/read_write.c:642	<code>loff_t pos = file_pos_read(f.file);</code>	
23	fs/read_write.c:612	<code>return file-&gt;f_pos;</code>	
24	fs/read_write.c:643	<code>ret = vfs_write(f.file, buf, count, &amp;pos);</code>	
25	fs/read_write.c:580	<code>if (!(file-&gt;f_mode &amp; FMODE_WRITE))</code>	
26	fs/read_write.c:582	<code>if (!(file-&gt;f_mode &amp; FMODE_CAN_WRITE))</code>	
27	fs/read_write.c:584	<code>if (unlikely(!access_ok(VERIFY_READ, buf, count)))</code>	
28	./arm/x86/include/asm/thread_info.h:162	<code>ti = (void *)(this_cpu_read_stable(kernel_stack) +</code>	
29	./arm/x86/include/asm/thread_info.h:163	<code>KERNEL_STACK_OFFSET - THREAD_SIZE);</code>	
30	./arm/x86/include/asm/thread_info.h:164	<code>return ti;</code>	
31	./arm/x86/include/asm/uaccess.h:57	<code>addr += size;</code>	
32	./arm/x86/include/asm/uaccess.h:58	<code>if (addr &lt; size)</code>	
33	./arm/x86/include/asm/uaccess.h:60	<code>return addr &gt; limit;</code>	
34	fs/read_write.c:587	<code>ret = rw_verify_area(WRITE, file, pos, count);</code>	
35	fs/read_write.c:391	<code>int retval = -EINVAL;</code>	

36	fs/read_write.c:393	inode = file_inode(file);
37	include/linux/fs.h:1963	<b>return f-&gt;f_inode;</b>
38	fs/read_write.c:394	<b>if</b> (unlikely((ssize_t) count < 0))
39	fs/read_write.c:396	pos = *ppos;
40	fs/read_write.c:397	<b>if</b> (unlikely(pos < 0)) {
41	fs/read_write.c:402	} <b>else if</b> (unlikely((loff_t) (pos + count) < 0)) {
42	fs/read_write.c:407	<b>if</b> (unlikely(inode->i_flctx && mandatory_lock(inode))) {
43	fs/read_write.c:414	retval = security_file_permission(file,
44	security/security.c:713	ret = security_ops->file_permission(file, mask);
45	security/selinux/hooks.c:3215	<b>struct</b> inode *inode = file_inode(file);
46	include/linux/fs.h:1963	<b>return f-&gt;f_inode;</b>
47	security/selinux/hooks.c:3216	<b>struct</b> file_security_struct *fsec = file->f_security;
48	security/selinux/hooks.c:3217	<b>struct</b> inode_security_struct *isec = inode->i_security;
49	security/selinux/hooks.c:218	<b>const struct</b> task_security_struct *tsec = current_security();
50	./arm/x86/include/asm/current.h:14	<b>return</b> this_cpu_read_stable(current_task);
51	security/selinux/hooks.c:218	<b>const struct</b> task_security_struct *tsec = current_security();
52	security/selinux/hooks.c:220	<b>return</b> tsec->sid;
53	security/selinux/hooks.c:3220	<b>if</b> (!mask)
54	security/selinux/avc.c:774	<b>return</b> avc_cache.latest_notif;
55	security/selinux/hooks.c:3227	<b>return</b> 0;
56	security/security.c:714	<b>if</b> (ret)
57	fs/read_write.c:635	SYSCALL_DEFINE3(write, <b>unsigned int</b> , fd, <b>const</b> <b>char</b> __user *, buf,

1~57 行會重複做 2-3 次不等，但是螢幕沒有印出第一行，猜測可能是有呼叫到 sys\_write 但是是把一些東西寫道 cache，並非螢幕上，例如：寫指令到 history

58	include/linux/fsnotify.h:40	<b>struct</b> path *path = &file->f_path;
59	include/linux/fs.h:1963	<b>return f-&gt;f_inode;</b>
60	include/linux/fsnotify.h:42	<b>_u32</b> fsnotify_mask = 0;
61	include/linux/fsnotify.h:45	<b>if</b> (file->f_mode & FMODE_NONOTIFY)
62	fs/read_write.c:416	<b>if</b> (retval)
63	fs/read_write.c:588	<b>if</b> (ret >= 0) {
64	fs/read_write.c:589	count = ret;
65	fs/read_write.c:590	file_start_write(file);
66	include/linux/fs.h:2381	<b>if</b> (!S_ISREG(file_inode(file)->i_mode))
67	include/linux/fs.h:1963	<b>return f-&gt;f_inode;</b>
68	include/linux/fs.h:2383	<b>_sb_start_write(file_inode(file)-&gt;i_sb,</b> <b>SB_FREEZE_WRITE</b> , true);

69	include/linux/fs.h:1963	<code>return f-&gt;f_inode;</code>	
70	include/linux/super.c:1194	<code>if (unlikely(sb-&gt;s_writers.frozen &gt;= level)) {</code>	
71	include/linux/super.c:1204	<code>percpu_counter_inc(&amp;sb-&gt;s_writers.counter[level-1]);</code>	
72	include/linux/percpu_counter.h:116	<code>percpu_counter_add(fbc, 1);</code>	
73	include/linux/percpu_counter.h:48	<code>__percpu_counter_add(fbc, amount,</code> <code>percpu_counter_batch);</code>	
74	include/linux/percpu_counter.:80	<code>count = __this_cpu_read(*fbc-&gt;counters) +</code> <code>amount;</code>	
75	include/linux/percpu-defs.:300	<code>static inline void __this_cpu_preempt_check(const</code> <code>char *op) { }</code>	
76	include/linux/percpu_counter.:81	<code>if (count &gt;= batch    count &lt;= -batch) {</code>	
77	include/linux/percpu_counter.:88	<code>this_cpu_add(*fbc-&gt;counters, amount);</code>	
78	include/linux/super.c:1204	<code>smp_mb();</code>	
79	include/linux/super.c:1210	<code>if (unlikely(sb-&gt;s_writers.frozen &gt;= level)) {</code>	
80	include/linux/super.c:1214	<code>return 1;</code>	
81	fs/read_write.c:591	<code>if (file-&gt;f_op-&gt;write)</code>	
82	fs/read_write.c:592	<code>ret = file-&gt;f_op-&gt;write(file, buf, count, pos);</code>	
83	fs/read_write.c:526	<code>struct iovec iov = { .iov_base = (void __user</code> <code>*)buf, .iov_len = len };</code>	
84	fs/read_write.c:531	<code>init_sync_kiocb(&amp;kiocb, filp);</code>	
85	./arm/x86/include/asm/aio.h:67	<code>.ki_obj.tsk = current,</code>	
86	./arm/x86/include/asm/current.h:14	<code>return this_cpu_read_stable(current_task);</code>	
87	./arm/x86/include/asm/aio.h:67	<code>.ki_obj.tsk = current,</code>	
88	./arm/x86/include/asm/aio.h:64	<code>*kiocb = (struct kiocb) {</code>	
89	fs/read_write.c:532	<code>kiocb.ki_pos = *ppos;</code>	
90	fs/read_write.c:533	<code>kiocb.ki_nbytes = len;</code>	
91	fs/read_write.c:534	<code>iov_iter_init(&amp;iter, WRITE, &amp;iov, 1, len);</code>	
92	Lib/iov_iter.c :325	<code>if (segment_eq(get_fs(), KERNEL_DS)) {</code>	
93	./arm/x86/include/asm/thread_info.h:162	<code>ti = (void *)(this_cpu_read_stable(kernel_stack) +</code>	
94	Lib/iov_iter.c :330	<code>i-&gt;type = direction;</code>	
95	fs/read_write.c:534	<code>ret = filp-&gt;f_op-&gt;write_iter(&amp;kiocb, &amp;iter);</code>	
96	mm/filemap.c:2656	<code>struct file *file = iocb-&gt;ki_filp;</code>	
97	Kernel/locking/mutex:97	<code>might_sleep();</code>	
98	Kernel/sched/core.c:4206	<code>if (should_resched()) {</code>	
99	./arm/x86/include/asm/preempt.h:95	<code>return</code> <code>unlikely(!raw_cpu_read_4(__preempt_count));</code>	
100	Kernel/sched/core.c:4206	<code>if (should_resched()) {</code>	
101	Kernel/sched/core.c:2845	<code>__preempt_count_add(PREEMPT_ACTIVE);</code>	
102	./arm/x86/include/asm/preempt.h:72	<code>raw_cpu_add_4(__preempt_count, val);</code>	

103	Kernel/sched/core.c:2845	<code>_schedule();</code>	
104	Kernel/sched/core.c:2721	<code>cpu = smp_processor_id();</code>	
105	Kernel/rcu/tree.c:285	<code>trace_rcu_utilization(TPS("Start context switch"));</code>	
… 中間開始進入要搶 lock 階段會經過 spinlock, ticklock, fair.c 等 function , 並且會常進出 tree.c 做資料結構處理，這段非常的冗長，詳情請見附錄(p37-p73)			
106	fs/read_write.c:536	<code>ret = filp-&gt;f_op-&gt;write_iter(&amp;kiocb, &amp;iter);</code>	
107	fs/read_write.c:537	<code>if (-EIOCBQUEUED == ret)</code>	
108	fs/read_write.c:537	<code>*ppos = kiocb.ki_pos;</code>	
109	fs/read_write.c:540	<code>return ret;</code>	
110	fs/read_write.c:597	<code>if (ret &gt; 0) {</code>	
111	fs/read_write.c:601	<code>inc_syscw(current);</code>	
112	fs/read_write.c:635	<code>SYSCALL_DEFINE3(write, unsigned int, fd, const char __user *, buf,</code>	

(4)印出第一行：

```
[ 17.414524] ls (902) used greatest stack depth: 12440 bytes left
/mnt/sharedFolder # cd hw03.syscall.asm/
/mnt/sharedFolder/hw03.syscall.asm # ls
makefile  syscall  syscall.c
/mnt/sharedFolder/hw03.syscall.asm # ./syscall
使用 'int 0x80' 呼叫 system call
```

## 附錄：

main.c process.c read\_write.c 0xffff0 "36

Name Type

fd un

buf co

count siz

f sti

ret ss

Console Registers Problems Executables Debugger Console Memory OS Resources Perf Profile View

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff81290945 in SYSC\_write+19 at fs/read\_write.c:638

SYSC\_write (fd=1, buf=0x2481390 "\nmount/sharedFolder/hw03.syscall.asm # /\033[J", count=1) at fs/read\_write.c:638

struct fd f = fdget\_pos(fd);

>>> █

main.c process.c read\_write.c 0xffff0 "36

EXPORT\_SYMBOL(vfs\_llseek);

static inline struct fd fdget\_pos(int fd)

{

return \_\_to\_fd(\_\_fdget\_pos(fd));

}

static inline void fdput\_pos(struct fd f)

{

if (f.flags & FDPUT\_POS\_UNLOCK)

mutex\_unlock(&f.file->f\_pos\_lock);

fdput(f);

}

SYSCALL\_DEFINE1(llseek, unsigned int, fd, off\_t, offset, unsigned

Console Registers Problems Executables Debugger Console

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff8128fa99 in fdget\_pos+11 at fs/read\_write.c:270

fdget\_pos (fd=1) at fs/read\_write.c:270

270 return \_\_to\_fd(\_\_fdget\_pos(fd));

>>> █

main.c process.c read\_write.c file.c 0xffff0 "35

unsigned long v = \_\_fdget(fd);

struct file \*file = (struct file \*) (v & ~3);

if (file && (file->f\_mode & FMODE\_ATOMIC\_POS)) {

if (file\_count(file) > 1) {

v |= FDPUT\_POS\_UNLOCK;

mutex\_lock(&file->f\_pos\_lock);

}

}

return v;

/\*

\* We only lock f\_pos if we have threads or if the file might be

\* shared with another process. In both cases we'll have an elevated

\* file count (done either by fdget() or by fork())

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff812be406 in \_\_fdget\_pos+11 at fs/file.c:708

\_\_fdget\_pos (fd=1) at fs/file.c:708

708 unsigned long v = \_\_fdget(fd);

>>> █

File: file.c Line: 697

```

692     return FDPUT_FPUT | (unsigned long)file;
693 }
694 }
695 unsigned long __fdget(unsigned int fd)
696 {
697     return __fget_light(fd, FMODE_PATH);
698 }
699 EXPORT_SYMBOL(__fdget);
700
701 unsigned long __fdget_raw(unsigned int fd)
702 {
703     return __fget_light(fd, 0);
704 }
705
706 unsigned long __fdget_pos(unsigned int fd)
707 {

```

Console [+] Threads [1] id 1 from 0xffffffff812be3ce in \_\_fdget+11 at fs/file.c:697

```

_fget(fd=1) at fs/file.c:697
697         return __fget_light(fd, FMODE_PATH);
>>> █

```

File: file.c Line: 680

```

675 * The fput_needed flag returned by fget_light should be passed to the
676 * corresponding fput_light.
677 */
678 static unsigned long __fget_light(unsigned int fd, fmode_t mask)
679 {
680     struct files_struct *files = current->files;
681     struct file *file;
682
683     if (atomic_read(&files->count) == 1) {
684         file = __fcheck_files(files, fd);
685         if (!file || unlikely(file->f_mode & mask))
686             return 0;
687         return (unsigned long)file;
688     } else {
689         file = __fget(fd, mask);
690         if (!file)

```

Console [+] Threads [1] id 1 from 0xffffffff812be32b in \_\_fget\_light+14 at fs/file.c:680

```

_fget_light (fd=1, mask=16384) at fs/file.c:680
680         struct files_struct *files = current->files;
>>> █

```

File: current.h Line: 14

```

1 #ifndef _ASSEMBLY_
2
3 struct task_struct;
4
5 DECLARE_PER_CPU(struct task_struct *, current_task);
6
7 static __always_inline struct task_struct *get_current(void)
8 {
8+14     return this_cpu_read_stable(current_task);
9 }
10
11 #define current get_current()
12
13 #endif /* __ASSEMBLY__ */
14
15 #endif /* _ASM_X86_CURRENT_H */
16
17 #endif /* _ASM_X86_CURRENT_H */
18
19 #endif /* _ASM_X86_CURRENT_H */
20
21 #endif /* _ASM_X86_CURRENT_H */
22

```

Console [+] Registers [+] Problems [+] Executables [+] Debugger Console [+] Memory [+] Threads [1] id 1 from 0xffffffff812be32b in get\_current+0 at ./arch/x86/include/asm/current.h:14

```

get_current () at ./arch/x86/include/asm/current.h:14
14         return this_cpu_read_stable(current_task);
>>> █

```

```

675 * The fput_needed flag returned by fget_light should be passed to the
676 * corresponding fput_light.
677 */
678 static unsigned long __fget_light(unsigned int fd, fmode_t mask)
679 {
680     struct files_struct *files = current->files;
681     struct file *file;
682
683     if (atomic_read(&files->count) == 1) {
684         file = __fcheck_files(files, fd);
685         if (!file || unlikely(file->f_mode & mask))
686             return 0;
687         return (unsigned long)file;
688     } else {
689         file = __fget(fd, mask);
690         if (!file)

```

Console [+] Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[1] id 1 from 0xffffffff812be33c in \_\_fget\_light+31 at fs/file.c:680

fget\_light (fd=1, mask=16384) at fs/file.c:680

680 struct files\_struct \*files = current->files;

>>>

```

675 * The fput_needed flag returned by fget_light should be passed to the
676 * corresponding fput_light.
677 */
678 static unsigned long __fget_light(unsigned int fd, fmode_t mask)
679 {
680     struct files_struct *files = current->files;
681     struct file *file;
682
683     if (atomic_read(&files->count) == 1) {
684         file = __fcheck_files(files, fd);
685         if (!file || unlikely(file->f_mode & mask))
686             return 0;
687         return (unsigned long)file;
688     } else {
689         file = __fget(fd, mask);
690         if (!file)

```

Console [+] Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg fd = 1

[1] id 1 from 0xffffffff812be347 in \_\_fget\_light+42 at fs/file.c:683

683 if (atomic\_read(&files->count) == 1) {

>>>

```

27     return ACCESS_ONCE((v)->counter);
28 }
29
30 /**
31  * atomic_set - set atomic variable
32  * @v: pointer of type atomic_t
33  * @i: required value
34  *
35  * Atomically sets the value of @v to @i.
36  */
37 static inline void atomic_set(atomic_t *v, int i)
38 {
39     v->counter = i;
40 }
41 /**
42 */

```

Console [+] Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[1] id 1 from 0xffffffff812bcba6 in atomic\_read+12 at ./arch/x86/include/asm/atomic.h

atomic\_read (v=0xffff88000ed06a00) at ./arch/x86/include/asm/atomic.h:27

27 return ACCESS\_ONCE((v)->counter);

>>>

```

675 * The fput_needed flag returned by fget_light should be passed to the
676 * corresponding fput_light.
677 */
678 static unsigned long __fget_light(unsigned int fd, fmode_t mask)
679 {
680     struct files_struct *files = current->files;
681     struct file *file;
682
683     if (atomic_read(&files->count) == 1) {
684         file = __fcheck_files(files, fd);
685         if (!file || unlikely(file->f_mode & mask))
686             return 0;
687         return (unsigned long)file;
688     } else {
689         file = __fget(fd, mask);
690         if (!file)

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff812be358 in \_\_fget\_light+59 at fs/file.c:684

```

__fget_light (fd=1, mask=16384) at fs/file.c:684
684             file = __fcheck_files(files, fd);
>>>
```

```

77 struct fdtable *fdt = rcu_dereference_raw(files->fdt);
78
79     if (fd < fdt->max_fds)
80         return rcu_dereference_raw(fdt->fd[fd]);
81     return NULL;
82 }
83
84 static inline struct file *fcheck_files(struct files_struct *files, unsig
85 {
86     rcu_lockdep_assert(rcu_read_lock_held() ||
87                         lockdep_is_held(&files->file_lock),
88                         "suspicious rcu_dereference_check() usage");
89     return __fcheck_files(files, fd);
90 }
91
92 */


```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff812bcf90 in \_\_fcheck\_files+15 at include/linux/fdtable.h

```

_fcheck_files (files=0xfffff88000ed06a00, fd=1) at include/linux/fdtable.h:77
77     struct fdtable *fdt = rcu_dereference_raw(files->fdt);
>>> 
```

```

77 struct fdtable *fdt = rcu_dereference_raw(files->fdt);
78
79     if (fd < fdt->max_fds)
80         return rcu_dereference_raw(fdt->fd[fd]);
81     return NULL;
82 }
83
84 static inline struct file *fcheck_files(struct files_struct *files, unsig
85 {
86     rcu_lockdep_assert(rcu_read_lock_held() ||
87                         lockdep_is_held(&files->file_lock),
88                         "suspicious rcu_dereference_check() usage");
89     return __fcheck_files(files, fd);
90 }
91
92 */


```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg mask = 16384

[+]

Threads

[1] id 1 from 0xffffffff812bcfb7 in \_\_fcheck\_files+54 at include/linux/fdtable.h

```

79     if (fd < fdt->max_fds)
>>>
```

```

77     struct fdtable *fdt = rcu_dereference_raw(files->fdt);
78
79     if (fd < fdt->max_fds)
80         return rcu_dereference_raw(fdt->fd[fd]);
81     return NULL;
82 }
83
84 static inline struct file *fcheck_files(struct files_struct *files, unsigned int fd)
85 {
86     rcu_lockdep_assert(rcu_read_lock_held() ||
87                         lockdep_is_held(&files->file_lock),
88                         "suspicious rcu_dereference_check() usage");
89     return __fcheck_files(files, fd);
90 }
91
92 */

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg mask = 16384

[+]

Threads

[1] id 1 from 0xffffffff812bcfc2 in \_\_fcheck\_files+65 at include/linux/fdtable.h:80

80 return rcu\_dereference\_raw(fdt->fd[fd]);

>>>

```

675 * The fput needed flag returned by fget_light should be passed to the
676 * corresponding fput_light.
677 */
678 static unsigned long __fget_light(unsigned int fd, fmode_t mask)
679 {
680     struct files_struct *files = current->files;
681     struct file *file;
682
683     if (atomic_read(&files->count) == 1) {
684         file = __fcheck_files(files, fd);
685         if (!file || unlikely(file->f_mode & mask))
686             return 0;
687         return (unsigned long)file;
688     } else {
689         file = fget(fd, mask);
690         if (!file)

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff812be36d in \_\_fget\_light+80 at fs/file.c:685

\_fget\_light (fd=1, mask=16384) at fs/file.c:685

685 if (!file || unlikely(file->f\_mode & mask))

>>>

```

677 */
678 static unsigned long __fget_light(unsigned int fd, fmode_t mask)
679 {
680     struct files_struct *files = current->files;
681     struct file *file;
682
683     if (atomic_read(&files->count) == 1) {
684         file = __fcheck_files(files, fd);
685         if (!file || unlikely(file->f_mode & mask))
686             return 0;
687         return (unsigned long)file;
688     } else {
689         file = fget(fd, mask);
690         if (!file)
691             return 0;
692         return FOPEN_FD | (unsigned long)file;

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg fd = 1

[+]

Threads

[1] id 1 from 0xffffffff812be392 in \_\_fget\_light+117 at fs/file.c:687

687 return (unsigned long)file;

>>>

```
704 }
705
706 unsigned long __fdget_pos(unsigned int fd)
707 {
708     unsigned long v = __fdget(fd);
709     struct file *file = (struct file *) (v & ~3);
710
711     if (file && (file->f_mode & FMODE_ATOMIC_POS)) {
712         if (file_count(file) > 1) {
713             v |= FDPUT_POS_UNLOCK;
714             mutex_lock(&file->f_pos_lock);
715         }
716     }
717     return v;
718 }
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff812be414 in \_\_fdget\_pos+25 at fs/file.c:709

fdget\_pos (fd=1) at fs/file.c:709

709 struct file \*file = (struct file \*) (v & ~3);

>>> █

```
704 }
705
706 unsigned long __fdget_pos(unsigned int fd)
707 {
708     unsigned long v = __fdget(fd);
709     struct file *file = (struct file *) (v & ~3);
710
711     if (file && (file->f_mode & FMODE_ATOMIC_POS)) {
712         if (file_count(file) > 1) {
713             v |= FDPUT_POS_UNLOCK;
714             mutex_lock(&file->f_pos_lock);
715         }
716     }
717     return v;
718 }
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg fd = 1

[+]

Threads

[1] id 1 from 0xffffffff812be420 in \_\_fdget\_pos+37 at fs/file.c:711

711 if (file && (file->f\_mode & FMODE\_ATOMIC\_POS)) {

>>> █

```
707 {
708     unsigned long v = __fdget(fd);
709     struct file *file = (struct file *) (v & ~3);
710
711     if (file && (file->f_mode & FMODE_ATOMIC_POS)) {
712         if (file_count(file) > 1) {
713             v |= FDPUT_POS_UNLOCK;
714             mutex_lock(&file->f_pos_lock);
715         }
716     }
717     return v;
718 }
719
720 /*
721 * We only lock f_pos if we have threads or if the file might be
722 * shared with another process. In both cases we'll have an elevated
723 */
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg fd = 1

[+]

Threads

[1] id 1 from 0xffffffff812be462 in \_\_fdget\_pos+103 at fs/file.c:717

717 return v;

>>> █

```

file.c current.h atomic.h fdtable.h file.h »35
50     return (struct fd){(struct file *) (v & ~3), v & 3};
51 }
52
53 static inline struct fd fdget(unsigned int fd)
54 {
55     return __to_fd(__fdget(fd));
56 }
57
58 static inline struct fd fdget_raw(unsigned int fd)
59 {
60     return __to_fd(__fdget_raw(fd));
61 }
62
63 extern int f_dupfd(unsigned int from, struct file *file, unsigned flags);
64 extern int replace_fd(unsigned fd, struct file *file, unsigned flags);
65 extern void set_close_on_exec(unsigned int fd, int flag).

```

Console Registers Problems Executables Debugger Console Memo

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8128f0b7 in \_\_to\_fd+12 at include/linux/file.h:50

to\_fd (v=1844661213255518464) at include/linux/file.h:50

50 return (struct fd){(struct file \*) (v & ~3), v & 3};

>>>

```

read_write.c file.c fdtable.h file.h »36
634
635 SYSCALL_DEFINE3(write, unsigned int, fd, const char __user *, buf,
636                 size_t, count)
637 {
638     struct fd f = fdget_pos(fd);
639     ssize_t ret = -EBADF;
640
641     if (f.file) {
642         loff_t pos = file_pos_read(f.file);
643         ret = vfs_write(f.file, buf, count, &pos);
644         if (ret >= 0)
645             file_pos_write(f.file, pos);
646         fdput_pos(f);
647     }
648
649     return ret.

```

Console Registers Problems Executables Debugger Console Memo

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81290957 in SYSC\_write+37 at fs/read\_write.c:639

SYSC\_write (fd=1, buf=0x2481390 "\nmnt/sharedFolder/hw03.syscall.asm # /\033[J'" 639 ssize\_t ret = -EBADF;

>>>

```

read_write.c file.c fdtable.h file.h »36
634
635 SYSCALL_DEFINE3(write, unsigned int, fd, const char __user *, buf,
636                 size_t, count)
637 {
638     struct fd f = fdget_pos(fd);
639     ssize_t ret = -EBADF;
640
641     if (f.file) {
642         loff_t pos = file_pos_read(f.file);
643         ret = vfs_write(f.file, buf, count, &pos);
644         if (ret >= 0)
645             file_pos_write(f.file, pos);
646         fdput_pos(f);
647     }
648
649     return ret.

```

Console Registers Problems Executables Debugger Console Memo

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg count = 1

[+]

Threads

[1] id 1 from 0xffffffff8129095f in SYSC\_write+45 at fs/read\_write.c:641

641 if (f.file) {

>>>

read\_write.c file.c fdtable.h file.h »36

```
634
635 SYSCALL_DEFINE3(write, unsigned int, fd, const char __user *, buf,
636             size_t, count)
637 {
638     struct fd f = fdget_pos(fd);
639     ssize_t ret = -EBADF;
640
641     if (f.file) {
642         loff_t pos = file_pos_read(f.file);
643         ret = vfs_write(f.file, buf, count, &pos);
644         if (ret >= 0)
645             file_pos_write(f.file, pos);
646         fdput_pos(f);
647     }
648
649     return ret;
650 }
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg count = 1

[+]

Threads

[1] id 1 from 0xffffffff81290968 in SYSC\_write+54 at fs/read\_write.c:642

```
642         loff_t pos = file_pos_read(f.file);
>>> █
```

read\_write.c file.c fdtable.h file.h »36

```
607 EXPORT_SYMBOL(vfs_write);
608
609 static inline loff_t file_pos_read(struct file *file)
610 {
611     return file->f_pos;
612 }
613
614 static inline void file_pos_write(struct file *file, loff_t pos)
615 {
616     file->f_pos = pos;
617 }
618
619
620 SYSCALL_DEFINE3(read, unsigned int, fd, char __user *, buf, size_t, count)
621 {
622     struct fd f = fdget_pos(fd);
623 }
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8129080d in file\_pos\_read+12 at fs/read\_write.c:612

```
file_pos_read (file=0xffff88000e61f200) at fs/read_write.c:612
612     return file->f_pos;
>>> █
```

read\_write.c file.c fdtable.h file.h »36

```
638     struct fd f = fdget_pos(fd);
639     ssize_t ret = -EBADF;
640
641     if (f.file) {
642         loff_t pos = file_pos_read(f.file);
643         ret = vfs_write(f.file, buf, count, &pos);
644         if (ret >= 0)
645             file_pos_write(f.file, pos);
646         fdput_pos(f);
647     }
648
649     return ret;
650 }
651
652 SYSCALL_DEFINE4(pread64, unsigned int, fd, char __user *, buf,
653                 size_t, count, loff_t, pos)
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81290978 in SYSC\_write+70 at fs/read\_write.c:643

```
SYSC_write (fd=1, buf=0x2481390 "\nmnt/sharedFolder/hw03.syscall.asm # /\033[J",
643             ret = vfs_write(f.file, buf, count, &pos);
>>> █
```

```
575
576 ssize_t vfs_write(struct file *file, const char __user *buf, size_t count)
577 {
578     ssize_t ret;
579
580     if (!(file->f_mode & FMODE_WRITE))
581         return -EBADF;
582     if (!(file->f_mode & FMODE_CAN_WRITE))
583         return -EINVAL;
584     if (unlikely(!access_ok(VERIFY_READ, buf, count)))
585         return -EFAULT;
586
587     ret = rw_verify_area(WRITE, file, pos, count);
588     if (ret >= 0) {
589         count = ret;
590         file_start_write(file);
591     }
592 }
```

```
Console Registers Problems Executables Debugger Console Memory
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff8129067c in vfs_write+24 at fs/read_write.c:580
vfs_write (file=0xffff88000e61f200, buf=0x2481390 "\nmnt/sharedFolder/hw03.syscal"
580         if (!(file->f_mode & FMODE_WRITE))
>>>
```

```
575
576 ssize_t vfs_write(struct file *file, const char __user *buf, size_t count)
577 {
578     ssize_t ret;
579
580     if (!(file->f_mode & FMODE_WRITE))
581         return -EBADF;
582     if (!(file->f_mode & FMODE_CAN_WRITE))
583         return -EINVAL;
584     if (unlikely(!access_ok(VERIFY_READ, buf, count)))
585         return -EFAULT;
586
587     ret = rw_verify_area(WRITE, file, pos, count);
588     if (ret >= 0) {
589         count = ret;
590         file_start_write(file);
591     }
592 }
```

```
Console Registers Problems Executables Debugger Console Memory
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
arg count = 1
[+]
Threads
[1] id 1 from 0xffffffff81290696 in vfs_write+50 at fs/read_write.c:582
582         if (!(file->f_mode & FMODE_CAN_WRITE))
>>> █
```

```
575
576 ssize_t vfs_write(struct file *file, const char __user *buf, size_t count)
577 {
578     ssize_t ret;
579
580     if (!(file->f_mode & FMODE_WRITE))
581         return -EBADF;
582     if (!(file->f_mode & FMODE_CAN_WRITE))
583         return -EINVAL;
584     if (unlikely(!access_ok(VERIFY_READ, buf, count)))
585         return -EFAULT;
586
587     ret = rw_verify_area(WRITE, file, pos, count);
588     if (ret >= 0) {
589         count = ret;
590         file_start_write(file);
591     }
592 }
```

```
Console Registers Problems Executables Debugger Console Memory
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
arg count = 1
[+]
Threads
[1] id 1 from 0xffffffff812906b2 in vfs_write+78 at fs/read_write.c:584
584         if (unlikely(!access_ok(VERIFY_READ, buf, count)))
>>>
```

```

read_write.c file.c file.h thread_info.h >36
162     ti = (void *) (this_cpu_read_stable(kernel_stack) +
163                     KERNEL_STACK_OFFSET - THREAD_SIZE);
164     return ti;
165 }
166
167 static inline unsigned long current_stack_pointer(void)
168 {
169     unsigned long sp;
170 #ifdef CONFIG_X86_64
171     asm("mov %%rsp,%0" : "=g" (sp));
172 #else
173     asm("mov %%esp,%0" : "=g" (sp));
174 #endif
175     return sp;
176 }
177

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8128efe5 in current\_thread\_info+8 at ./arch/x86/include/as

current\_thread\_info () at ./arch/x86/include/asm/thread\_info.h:162

162 ti = (void \*) (this\_cpu\_read\_stable(kernel\_stack) +

>>>

```

read_write.c file.c file.h thread_info.h >36
162     ti = (void *) (this_cpu_read_stable(kernel_stack) +
163                     KERNEL_STACK_OFFSET - THREAD_SIZE);
164     return ti;
165 }
166
167 static inline unsigned long current_stack_pointer(void)
168 {
169     unsigned long sp;
170 #ifdef CONFIG_X86_64
171     asm("mov %%rsp,%0" : "=g" (sp));
172 #else
173     asm("mov %%esp,%0" : "=g" (sp));
174 #endif
175     return sp;
176 }
177

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg pos = 0xfffff88000e5bff18: 0

[+]

Threads

[1] id 1 from 0xffffffff8128eff7 in current\_thread\_info+26 at ./arch/x86/include

163 KERNEL\_STACK\_OFFSET - THREAD\_SIZE);

>>>

```

read_write.c file.c file.h thread_info.h >36
162     ti = (void *) (this_cpu_read_stable(kernel_stack) +
163                     KERNEL_STACK_OFFSET - THREAD_SIZE);
164     return ti;
165 }
166
167 static inline unsigned long current_stack_pointer(void)
168 {
169     unsigned long sp;
170 #ifdef CONFIG_X86_64
171     asm("mov %%rsp,%0" : "=g" (sp));
172 #else
173     asm("mov %%esp,%0" : "=g" (sp));
174 #endif
175     return sp;
176 }
177

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8128efe5 in current\_thread\_info+8 at ./arch/x86/include/as

current\_thread\_info () at ./arch/x86/include/asm/thread\_info.h:162

162 ti = (void \*) (this\_cpu\_read\_stable(kernel\_stack) +

>>>

```

162     ti = (void *)this_cpu_read_stable(kernel_stack) +
163             KERNEL_STACK_OFFSET - THREAD_SIZE);
164     return ti;
165 }
166
167 static inline unsigned long current_stack_pointer(void)
168 {
169     unsigned long sp;
170 #ifdef CONFIG_X86_64
171     asm("mov %%rsp,%0" : "=g" (sp));
172 #else
173     asm("mov %%esp,%0" : "=g" (sp));
174 #endif
175     return sp;
176 }
177

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg pos = 0xfffff88000e5bff18: 0

[+]

Threads

[1] id 1 from 0xffffffff8128f001 in current\_thread\_info+36 at ./arch/x86/include/as

164 return ti;

>>> s

```

57     addr += size;
58     if (addr < size)
59         return true;
60     return addr > limit;
61 }
62
63 #define __range_not_ok(addr, size, limit) \
64 ({ \
65     __chk_user_ptr(addr); \
66     __chk_range_not_ok((unsigned long __force)(addr), size, limit); \
67 })
68
69 /**
70 * access_ok: - Checks if a user space pointer is valid
71 * @type: Type of access: %VERIFY_READ or %VERIFY_WRITE. Note that
72 * %VRFY_WRTTF is a superset of %VRFY_RAN - if it is safe

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8128f502 in \_\_chk\_range\_not\_ok+20 at ./arch/x86/include/as

\_\_chk\_range\_not\_ok (addr=38278032, size=1, limit=140737488351232) at ./arch/x86/include/as

57 addr += size;

>>> s

```

57     addr += size;
58     if (addr < size)
59         return true;
60     return addr > limit;
61 }
62
63 #define __range_not_ok(addr, size, limit) \
64 ({ \
65     __chk_user_ptr(addr); \
66     __chk_range_not_ok((unsigned long __force)(addr), size, limit); \
67 })
68
69 /**
70 * access_ok: - Checks if a user space pointer is valid
71 * @type: Type of access: %VERIFY_READ or %VERIFY_WRITE. Note that
72 * %VRFY_WRTTF is a superset of %VRFY_RAN - if it is safe

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg pos = 0xfffff88000e5bff18: 0

[+]

Threads

[1] id 1 from 0xffffffff8128f50a in \_\_chk\_range\_not\_ok+28 at ./arch/x86/include/as

58 if (addr < size)

>>> s

```

read_write.c file.h thread_info.h uaccess.h >36
57     addr += size;
58     if (addr < size)
59         return true;
60     return addr > limit;
61 }
62
63 #define __range_not_ok(addr, size, limit) \
64 ({ \
65     __chk_user_ptr(addr); \
66     __chk_range_not_ok((unsigned long __force)(addr), size, limit); \
67 })
68
69 /**
70 * access_ok: - Checks if a user space pointer is valid
71 * @type: Type of access: %VERIFY_READ or %VERIFY_WRITE. Note that
72 *       %VERIFY_WRTTF is a superset of %VERIFY_READ - if it is safe

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg pos = 0xfffff88000e5bff18: 0

[+]

Threads

[1] id 1 from 0xffffffff8128f51b in \_\_chk\_range\_not\_ok+45 at ./arch/x86/include,

60 return addr > limit;

>>> █

```

read_write.c file.h thread_info.h uaccess.h >36
577 {
578     ssize_t ret;
579
580     if (!(file->f_mode & FMODE_WRITE))
581         return -EBADF;
582     if (!(file->f_mode & FMODE_CAN_WRITE))
583         return -EINVAL;
584     if (unlikely(!access_ok(VERIFY_READ, buf, count)))
585         return -EFAULT;
586
587     ret = rw_verify_area(WRITE, file, pos, count);
588     if (ret >= 0) {
589         count = ret;
590         file_start_write(file);
591         if (file->f_op->write)
592             ret = file->f_op->write(file, buf, count, nns);

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff812906e5 in vfs\_write+129 at fs/read\_write.c:587

vfs\_write (file=0xfffff88000e61f200, buf=0x2481390 "\nmnt/sharedFolder/hw03.syscall.a

587 ret = rw\_verify\_area(WRITE, file, pos, count);

>>> █

```

read_write.c file.h thread_info.h uaccess.h >36
386 */
387 int rw_verify_area(int read_write, struct file *file, const loff_t *ppos,
388 {
389     struct inode *inode;
390     loff_t pos;
391     int retval = -EINVAL;
392
393     inode = file_inode(file);
394     if (unlikely((ssize_t) count < 0))
395         return retval;
396     pos = *ppos;
397     if (unlikely(pos < 0)) {
398         if (!unsigned_offsets(file))
399             return retval;
400         if (count >= -pos) /* both values are in 0..LLONG_MAX */
401             return -EOVFLW;

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8128feb in rw\_verify\_area+23 at fs/read\_write.c:391

rw\_verify\_area (read\_write=1, file=0xfffff88000e61f200, ppos=0xfffff88000e5bff18, c

391 int retval = -EINVAL;

>>> █

```
386 */
387 int rw_verify_area(int read_write, struct file *file, const loff_t *ppos,
388 {
389     struct inode *inode;
390     loff_t pos;
391     int retval = -EINVAL;
392
393     inode = file_inode(file);
394     if (unlikely((ssize_t) count < 0))
395         return retval;
396     pos = *ppos;
397     if (unlikely(pos < 0)) {
398         if (!unsigned_offsets(file))
399             return retval;
400         if (count >= -pos) /* both values are in 0..LLONG_MAX */
401             return -EINVAL;
402
403     inode = file_inode(file);
404 }
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff8128fec2 in rw\_verify\_area+30 at fs/read\_write.c:393

393 inode = file\_inode(file);

>>>

```
1963     return f->f_inode;
1964 }
1965
1966 /* /sys/fs */
1967 extern struct kobject *fs_kobj;
1968
1969 #define MAX_RW_COUNT (INT_MAX & PAGE_CACHE_MASK)
1970
1971 #define FLOCK_VERIFY_READ 1
1972 #define FLOCK_VERIFY_WRITE 2
1973
1974 #ifdef CONFIG_FILE_LOCKING
1975 extern int locks_mandatory_locked(struct file *);
1976 extern int locks_mandatory_area(int, struct inode *, struct file *, loff_t, size_t, int);
1977 */
1978 /*
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff8128f18d in file\_inode+12 at include/linux/fs.h:1963

file\_inode (f=0xffff88000e61f200) at include/linux/fs.h:1963

1963 return f->f\_inode;

>>>

```
386 */  
387 int rw_verify_area(int read_write, struct file *file, const loff_t *ppos  
388 {  
389     struct inode *inode;  
390     loff_t pos;  
391     int retval = -EINVAL;  
392  
393     inode = file_inode(file);  
394     if (unlikely((ssize_t) count < 0))  
395         return retval;  
396     pos = *ppos;  
397     if (unlikely(pos < 0)) {  
398         if (!unsigned_offsets(file))  
399             return retval;  
400         if (count >= -pos) /* both values are in 0..LLONG_MAX */  
401             return -EOVERFLOW;
```

```
Console Registers Problems Executables Debugger Console Memory  
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)  
[+]  
Threads  
[1] id 1 from 0xffffffff8128fed2 in rw_verify_area+46 at fs/read_write.c:394  
rw_verify_area (read_write=1, file=0xffff88000e61f200, ppos=0xffff88000e5bff18,  
394         if (unlikely((ssize_t) count < 0))  
>>> █
```

```
386 */  
387 int rw_verify_area(int read_write, struct file *file, const loff_t *ppos  
388 {  
389     struct inode *inode;  
390     loff_t pos;  
391     int retval = -EINVAL;  
392  
393     inode = file_inode(file);  
394     if (unlikely((ssize_t) count < 0))  
395         return retval;  
396     pos = *ppos;  
397     if (unlikely(pos < 0)) {  
398         if (!unsigned_offsets(file))  
399             return retval;  
400         if (count >= -pos) /* both values are in 0..LLONG_MAX */  
401             return -EOVERFLOW;
```

```
Console Registers Problems Executables Debugger Console Memory  
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)  
arg pos = 0xffff88000e5bff18: 0  
[+]  
Threads  
[1] id 1 from 0xffffffff8128fea in rw_verify_area+70 at fs/read_write.c:396  
396         pos = *ppos;  
>>> █
```

```
387 int rw_verify_area(int read_write, struct file *file, const loff_t *ppos  
388 {  
389     struct inode *inode;  
390     loff_t pos;  
391     int retval = -EINVAL;  
392  
393     inode = file_inode(file);  
394     if (unlikely((ssize_t) count < 0))  
395         return retval;  
396     pos = *ppos;  
397     if (unlikely(pos < 0)) {  
398         if (!unsigned_offsets(file))  
399             return retval;  
400         if (count >= -pos) /* both values are in 0..LLONG_MAX */  
401             return -EOVERFLOW;  
402     } else if (unlikely((loff_t) (pos + count) < 0)) {
```

```
Console Registers Problems Executables Debugger Console Memory  
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)  
arg pos = 0xffff88000e5bff18: 0  
[+]  
Threads  
[1] id 1 from 0xffffffff8128fef5 in rw_verify_area+81 at fs/read_write.c:397  
397         if (unlikely(pos < 0)) {  
>>> █
```

```

392     inode = file_inode(file);
393     if (unlikely((ssize_t) count < 0))
394         return retval;
395     pos = *ppos;
396     if (unlikely(pos < 0)) {
397         if (!unsigned_offsets(file))
398             return retval;
399         if (count >= -pos) /* both values are in 0..LLONG_MAX */
400             return -EOVERFLOW;
401     } else if (unlikely((loff_t) (pos + count) < 0)) {
402         if (!unsigned_offsets(file))
403             return retval;
404     }
405
406     if (unlikely(inode->i_fctx && mandatory_lock(inode))) {
407

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg pos = 0xfffff88000e5bff18: 0

[+]

Threads

[1] id 1 from 0xffffffff8128ff34 in rw\_verify\_area+144 at fs/read\_write.c:407

402 } else if (unlikely((loff\_t) (pos + count) < 0)) {  
>>>

```

397     if (unlikely(pos < 0)) {
398         if (!unsigned_offsets(file))
399             return retval;
400         if (count >= -pos) /* both values are in 0..LLONG_MAX */
401             return -EOVERFLOW;
402     } else if (unlikely((loff_t) (pos + count) < 0)) {
403         if (!unsigned_offsets(file))
404             return retval;
405     }
406
407     if (unlikely(inode->i_fctx && mandatory_lock(inode))) {
408         retval = locks_mandatory_area(
409             read_write == READ ? FLOCK_VERIFY_READ : FLOCK_VERIFY_WRITE,
410             inode, file, pos, count);
411         if (retval < 0)
412             return retval.
413

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg pos = 0xfffff88000e5bff18: 0

[+]

Threads

[1] id 1 from 0xffffffff8128ff63 in rw\_verify\_area+191 at fs/read\_write.c:407

407 if (unlikely(inode->i\_fctx && mandatory\_lock(inode))) {  
>>>

```

409         read_write == READ ? FLOCK_VERIFY_READ : FLOCK_VERIFY_WRITE,
410         inode, file, pos, count);
411     if (retval < 0)
412         return retval;
413 }
414     if (read_write == READ ? MAY_READ : MAY_WRITE);
415     if (retval)
416         return retval;
417     return count > MAX_RW_COUNT ? MAX_RW_COUNT : count;
418 }
419
420 ssize_t do_sync_read(struct file *filp, char __user *buf, size_t len, lo
421 {
422     struct iovec iov = { .iov_base = buf, .iov_len = len };
423     struct kinch kinch;
424

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg pos = 0xfffff88000e5bff18: 0

[+]

Threads

[1] id 1 from 0xffffffff8128ffcf in rw\_verify\_area+299 at fs/read\_write.c:414

414 retval = security\_file\_permission(file,  
>>>

```

read_write.c fs.h uaccess.h security.c "36"
713     ret = security_ops->file_permission(file, mask);
714     if (ret)
715         return ret;
716
717     return fsnotify_perm(file, mask);
718 }
719
720 int security_file_alloc(struct file *file)
721 {
722     return security_ops->file_alloc_security(file);
723 }
724
725 void security_file_free(struct file *file)
726 {
727     security_ops->file_free_security(file);
728 }

Console Registers Problems Executables Debugger Console Memory
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff814aae88 in security_file_permission+15 at security/security
security_file_permission (file=0xffff88000e61f200, mask=2) at security/security.c:71
713     ret = security_ops->file_permission(file, mask);
>>>

```

```

read_write.c fs.h security.c hooks.c "36"
3215 struct inode *inode = file_inode(file);
3216 struct file_security_struct *fsec = file->f_security;
3217 struct inode_security_struct *isec = inode->i_security;
3218 u32 sid = current_sid();
3219
3220 if (!mask)
3221     /* No permission to check. Existence test. */
3222     return 0;
3223
3224 if (sid == fsec->sid && fsec->isid == isec->sid &&
3225     fsec->pseqno == avc_policy_seqno())
3226     /* No change since file_open check. */
3227     return 0;
3228
3229 return selinux_revalidate_file_permission(file, mask);
3230 }

Console Registers Problems Executables Debugger Console Memory
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff814b5194 in selinux_file_permission+16 at security/seli
selinux_file_permission (file=0xffff88000e61f200, mask=2) at security/selinux.h:3215
3215     struct inode *inode = file_inode(file);
>>>

```

```

read_write.c fs.h security.c hooks.c "36"
1963     return f->f_inode;
1964 }
1965
1966 /* /sys/fs */
1967 extern struct kobject *fs_kobj;
1968
41 1969 #define MAX_RW_COUNT (INT_MAX & PAGE_CACHE_MASK)
1970
1971 #define FLOCK_VERIFY_READ 1
1972 #define FLOCK_VERIFY_WRITE 2
1973
1974 #ifdef CONFIG_FILE_LOCKING
1975 extern int locks_mandatory_locked(struct file *);
1976 extern int locks_mandatory_area(int, struct inode *, struct file *, loff_t,
1977 /*

Console Registers Problems Executables Debugger Console Memory
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff814af336 in file_inode+12 at include/linux/fs.h:1963
file_inode (f=0xffff88000e61f200) at include/linux/fs.h:1963
1963     return f->f_inode;
>>>

```

```

3215     struct inode *inode = file_inode(file);
3216     struct file_security_struct *fsec = file->f_security;
3217     struct inode_security_struct *isec = inode->i_security;
3218     u32 sid = current_sid();
3219
3220     if (!mask)
3221         /* No permission to check. Existence test. */
3222         return 0;
3223
3224     if (sid == fsec->sid && fsec->isid == isec->sid &&
3225         fsec->pseqno == avc_policy_seqno())
3226         /* No change since file_open check. */
3227         return 0;
3228
3229     return selinux_revalidate_file_permission(file, mask);
3230 }

```

Console Registers Problems Executables Debugger Console Memo

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff814b51a4 in **selinux\_file\_permission+32** at **security/selinux/file\_permission** (file=0xffff88000e61f200, mask=2) at security/selinux/file\_permission

3216 struct file\_security\_struct \*fsec = file->f\_security;

>>>

```

3215     struct inode *inode = file_inode(file);
3216     struct file_security_struct *fsec = file->f_security;
3217     struct inode_security_struct *isec = inode->i_security;
3218     u32 sid = current_sid();
3219
3220     if (!mask)
3221         /* No permission to check. Existence test. */
3222         return 0;
3223
3224     if (sid == fsec->sid && fsec->isid == isec->sid &&
3225         fsec->pseqno == avc_policy_seqno())
3226         /* No change since file_open check. */
3227         return 0;
3228
3229     return selinux_revalidate_file_permission(file, mask);
3230 }

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg mask = 2

[+]

Threads

[1] id 1 from 0xffffffff814b51b3 in **selinux\_file\_permission+47** at **security/selinux/file\_permission**

3217 struct inode\_security\_struct \*isec = inode->i\_security;

>>>

```

213 /*
214  * get the subjective security ID of the current task
215  */
216 static inline u32 current_sid(void)
217 {
218     const struct task_security_struct *tsec = current_security();
219
220     return tsec->sid;
221 }
222
223 /* Allocate and free functions for each kind of security blob. */
224
225 static int inode_alloc_security(struct inode *inode)
226 {
227     struct inode_security_struct *isec;
228     u32 sid = current_sid();

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff814afa51 in **current\_sid+8** at **security/selinux/hooks.c:218**

current\_sid () at security/selinux/hooks.c:218

218 const struct task\_security\_struct \*tsec = current\_security();

>>>

fs.h current.h security.c hooks.c »38

```

1 #ifndef __ASSEMBLY__
2 struct task_struct;
3
4 DECLARE_PER_CPU(struct task_struct *, current_task);
5
6 static __always_inline struct task_struct *get_current(void)
7 {
8     return this_cpu_read_stable(current_task);
9 }
10
11 #define current get_current()
12
13 #endif /* __ASSEMBLY__ */
14
15 #endif /* _ASM_X86_CURRENT_H */

```

Console Registers Problems Executables Debugger Console Memory OS R

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff814afa51 in get\_current+0 at ./arch/x86/include/asm/current.h:14

get\_current () at ./arch/x86/include/asm/current.h:14

14 return this\_cpu\_read\_stable(current\_task);

>>>

fs.h current.h security.c hooks.c »38

```

213 /*
214  * get the subjective security ID of the current task
215  */
216 static inline u32 current_sid(void)
217 {
218     const struct task_security_struct *tsec = current_security();
219
220     return tsec->sid;
221 }
222
223 /* Allocate and free functions for each kind of security blob. */
224
225 static int inode_alloc_security(struct inode *inode)
226 {
227     struct inode_security_struct *isec;
228     u32 sid = current_sid();

```

Console Registers Problems Executables Debugger Console Memory OS R

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff814afa62 in current\_sid+25 at security/selinux/hooks.c:218

current\_sid () at security/selinux/hooks.c:218

218 const struct task\_security\_struct \*tsec = current\_security();

>>>

fs.h current.h security.c hooks.c »38

```

213 /*
214  * get the subjective security ID of the current task
215  */
216 static inline u32 current_sid(void)
217 {
218     const struct task_security_struct *tsec = current_security();
219
220     return tsec->sid;
221 }
222
223 /* Allocate and free functions for each kind of security blob. */
224
225 static int inode_alloc_security(struct inode *inode)
226 {
227     struct inode_security_struct *isec;
228     u32 sid = current_sid();

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg mask = 2

[+]

Threads

[1] id 1 from 0xffffffff814afa71 in current\_sid+40 at security/selinux/hooks.c:220

220 return tsec->sid;

>>>

```

3215     struct inode *inode = file_inode(file);
3216     struct file_security_struct *fsec = file->f_security;
3217     struct inode_security_struct *isec = inode->i_security;
3218     u32 sid = current_sid();
3219
3220     if (!mask)
3221         /* No permission to check. Existence test. */
3222         return 0;
3223
3224     if (sid == fsec->sid && fsec->isid == isec->sid &&
3225         fsec->pseqno == avc_policy_seqno())
3226         /* No change since file_open check. */
3227         return 0;
3228
3229     return selinux_revalidate_file_permission(file, mask);
3230 }

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff814b51c7 in **selinux\_file\_permission+67** at **security/se**

selinux\_file\_permission (file=0xffff88000e61f700, mask=2) at **security/selinux/se**

3220 if (!mask)

>>>

```

774     return avc_cache.latest_notif;
775 }
776
777 void avc_disable(void)
778 {
779     /*
780     * If you are looking at this because you have realized that we are
781     * not destroying the avc_node_cachep it might be easy to fix, but
782     * I don't know the memory barrier semantics well enough to know. It
783     * possible that some other task dereferenced security_ops when
784     * it still pointed to selinux operations. If that is the case it's
785     * possible that it is about to use the avc and is about to need the
786     * avc_node_cachep. I know I could wrap the security.c security ops
787     * in an rcu_lock, but seriously, it's not worth it. Instead I just
788     * the cache and get that memory back.
789     */

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff814aead4 in **avc\_policy\_seqno+4** at **security/selinux/avc.c:**

avc\_policy\_seqno () at **security/selinux/avc.c:774**

774 return avc\_cache.latest\_notif;

>>>

```

3217     struct inode_security_struct *isec = inode->i_security;
3218     u32 sid = current_sid();
3219
3220     if (!mask)
3221         /* No permission to check. Existence test. */
3222         return 0;
3223
3224     if (sid == fsec->sid && fsec->isid == isec->sid &&
3225         fsec->pseqno == avc_policy_seqno())
3226         /* No change since file_open check. */
3227         return 0;
3228
3229     return selinux_revalidate_file_permission(file, mask);
3230 }
3231
3232 static int selinux_file_alloc_security(struct file *file)

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg mask = 2

[+]

Threads

[1] id 1 from 0xffffffff814b5201 in **selinux\_file\_permission+125** at **security/se**

3227 return 0;

>>>

```

713     ret = security_ops->file_permission(file, mask);
714     if (ret)
715         return ret;
716
717     return fsnotify_perm(file, mask);
718 }
719
720 int security_file_alloc(struct file *file)
721 {
722     return security_ops->file_alloc_security(file);
723 }
724
725 void security_file_free(struct file *file)
726 {
727     security_ops->file_free_security(file);
728 }

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff814aaea7 in security\_file\_permission+46 at security/security.c:71

security\_file\_permission (file=0xffff88000e61f700, mask=2) at security/security.c:71

714 if (ret)

>>>

```

630     fdput_pos(f);
631 }
632     return ret;
633 }
634
635 SYSCALL_DEFINE3(write, unsigned int, fd, const char __user *, buf,
636                 size_t, count)
637 {
638     struct fd f = fdget_pos(fd);
639     ssize_t ret = -EBADF;
640
641     if (f.file) {
642         loff_t pos = file_pos_read(f.file);
643         ret = vfs_write(f.file, buf, count, &pos);
644         if (ret >= 0)
645             file_nos_write(f.file_nos).

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff81290912 in Sys\_write+20 at fs/read\_write.c:635

Breakpoint 1, Sys\_write (fd=3, buf=38271712, count=10) at fs/read\_write.c:635

635 SYSCALL\_DEFINE3(write, unsigned int, fd, const char \_\_user \*, buf,

>>>

```

40     struct path *path = &file->f_path;
41     struct inode *inode = file_inode(file);
42     __u32 fsnotify_mask = 0;
43     int ret;
44
45     if (file->f_mode & FMODE_NONOTIFY)
46         return 0;
47     if (!(mask & (MAY_READ | MAY_OPEN)))
48         return 0;
49     if (mask & MAY_OPEN)
50         fsnotify_mask = FS_OPEN_PERM;
51     else if (mask & MAY_READ)
52         fsnotify_mask = FS_ACCESS_PERM;
53     else
54         BUG();
55

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff814a9c7b in fsnotify\_perm+15 at include/linux/fsnotify.h

fsnotify\_perm (file=0xffff88000e61f700, mask=2) at include/linux/fsnotify.h:40

40 struct path \*path = &file->f\_path;

>>>

```

1963     return f->f_inode;
1964 }
1965
1966 /* /sys/fs */
1967 extern struct kobject *fs_kobj;
1968
1969 #define MAX_RW_COUNT (INT_MAX & PAGE_CACHE_MASK)
1970
1971 #define FLOCK_VERIFY_READ 1
1972 #define FLOCK_VERIFY_WRITE 2
1973
1974 #ifdef CONFIG_FILE_LOCKING
1975 extern int locks_mandatory_locked(struct file *);
1976 extern int locks_mandatory_area(int, struct inode *, struct file *, loff_t);
1977 */
1978 */

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff814a9ad5 in file\_inode+12 at include/linux/fs.h:1963

file\_inode (f=0xffff88000e61f700) at include/linux/fs.h:1963

1963 return f->f\_inode;

>>>

```

40     struct path *path = &file->f_path;
41     struct inode *inode = file_inode(file);
42     u32 fsnotify_mask = 0;
43     int ret;
44
45     if (file->f_mode & FMODE_NONOTIFY)
46         return 0;
47     if (!(mask & (MAY_READ | MAY_OPEN)))
48         return 0;
49     if (mask & MAY_OPEN)
50         fsnotify_mask = FS_OPEN_PERM;
51     else if (mask & MAY_READ)
52         fsnotify_mask = FS_ACCESS_PERM;
53     else
54         BUG();
55

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff814a9c97 in fsnotify\_perm+43 at include/linux/fsnotify.h:4

fsnotify\_perm (file=0xffff88000e61f700, mask=2) at include/linux/fsnotify.h:42

42 \_u32 fsnotify\_mask = 0;

>>>

```

40     struct path *path = &file->f_path;
41     struct inode *inode = file_inode(file);
42     u32 fsnotify_mask = 0;
43     int ret;
44
45     if (file->f_mode & FMODE_NONOTIFY)
46         return 0;
47     if (!(mask & (MAY_READ | MAY_OPEN)))
48         return 0;
49     if (mask & MAY_OPEN)
50         fsnotify_mask = FS_OPEN_PERM;
51     else if (mask & MAY_READ)
52         fsnotify_mask = FS_ACCESS_PERM;
53     else
54         BUG();
55

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg mask = 2

[+]

Threads

[1] id 1 from 0xffffffff814a9c9e in fsnotify\_perm+50 at include/linux/fsnotify.h:45

45 if (file->f\_mode & FMODE\_NONOTIFY)

>>>

```

409         read_write == READ ? FLOCK_VERIFY_READ : FLOCK_VERIFY_WRITE,
410         inode, file, pos, count);
411     if (retval < 0)
412         return retval;
413 }
414     retval = security_file_permission(file,
415         read_write == READ ? MAY_READ : MAY_WRITE);
416     if (retval)
417         return retval;
418     return count > MAX_RW_COUNT ? MAX_RW_COUNT : count;
419 }
420
421 ssize_t do_sync_read(struct file *filp, char __user *buf, size_t len, lo
422 {
423     struct iovec iov = { .iov_base = buf, .iov_len = len };
424     struct kiocb kiocb;

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8128ffff in rw\_verify\_area+334 at fs/read\_write.c:416

rw\_verify\_area (read\_write=1, file=0xffff88000e61f700, ppos=0xffff88000e5bfff18, co

416 if (retval)

>>>

```

583     return -EINVAL;
584     if (unlikely(!access_ok(VIEWER_READ, buf, count)))
585         return -EFAULT;
586
587     ret = rw_verify_area(WRITE, file, pos, count);
588     if (ret >= 0) {
589         count = ret;
590         file_start_write(file);
591         if (file->f_op->write)
592             ret = file->f_op->write(file, buf, count, pos);
593         else if (file->f_op->aio_write)
594             ret = do_sync_write(file, buf, count, pos);
595         else
596             ret = new_sync_write(file, buf, count, pos);
597         if (ret > 0) {
598             fsnotify_mkdir(file);

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81290704 in vfs\_write+160 at fs/read\_write.c:588

vfs\_write (file=0xffff88000e61f700, buf=0x247fae0 "./syscall\nolder/hw03.syscall.asm"
588 if (ret >= 0) {

>>>

```

583     return -EINVAL;
584     if (unlikely(!access_ok(VIEWER_READ, buf, count)))
585         return -EFAULT;
586
587     ret = rw_verify_area(WRITE, file, pos, count);
588     if (ret >= 0) {
589         count = ret;
590         file_start_write(file);
591         if (file->f_op->write)
592             ret = file->f_op->write(file, buf, count, pos);
593         else if (file->f_op->aio_write)
594             ret = do_sync_write(file, buf, count, pos);
595         else
596             ret = new_sync_write(file, buf, count, pos);
597         if (ret > 0) {
598             fsnotify_mkdir(file);

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg count = 10

[+]

Threads

[1] id 1 from 0xffffffff8129070f in vfs\_write+171 at fs/read\_write.c:589

589 count = ret;

>>>

read\_write.c

```

583     return -EINVAL;
584     if (unlikely(!access_ok(VERIFY_READ, buf, count)))
585         return -EFAULT;
586
587     ret = rw_verify_area(WRITE, file, pos, count);
588     if (ret >= 0) {
589         count = ret;
590         file_start_write(file);
591         if (file->f_op->write)
592             ret = file->f_op->write(file, buf, count, pos);
593         else if (file->f_op->aio write)
594             ret = do_sync_write(file, buf, count, pos);
595         else
596             ret = new_sync_write(file, buf, count, pos);
597         if (ret > 0)
598             fnntifv mndifv(file).

```

Console

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
arg count = 10
[+]
Threads
[1] id 1 from 0xffffffff81290717 in vfs_write+179 at fs/read_write.c:590
590         file_start_write(file);
>>>

```

read\_write.c

```

2376     return (inode->i_mode & S_IXUGO) || S_ISDIR(inode->i_mode);
2377 }
2378
2379 static inline void file_start_write(struct file *file)
2380 {
2381     if (!S_ISREG(file_inode(file)->i_mode))
2382         return;
2383     __sb_start_write(file_inode(file)->i_sb, SB_FREEZE_WRITE, true);
2384 }
2385
2386 static inline bool file_start_write_trylock(struct file *file)
2387 {
2388     if (!S_ISREG(file_inode(file)->i_mode))
2389         return true;
2390     return __sb_start_write(file_inode(file)->i_sb, SB_FREEZE_WRITE, fal
2391 }

```

Console

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff8128f209 in file_start_write+12 at include/linux/fs.h:2381
file_start_write (file=0xffff88000e61f700) at include/linux/fs.h:2381
2381     if (!S_ISREG(file_inode(file)->i_mode))
>>>

```

read\_write.c

```

1958 extern void iput(struct inode *);
1959 extern int generic_update_time(struct inode *, struct timespec *, int);
1960
1961 static inline struct inode *file_inode(const struct file *f)
1962 {
1963     return f->f_inode;
1964 }
1965
1966 /* /sys/fs */
1967 extern struct kobject *fs_kobj;
1968
1969 #define MAX_RW_COUNT (INT_MAX & PAGE_CACHE_MASK)
1970
1971 #define FLOCK_VERIFY_READ 1
1972 #define FLOCK_VERIFY_WRITE 2
1973

```

Console

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff8128f18d in file_inode+12 at include/linux/fs.h:1963
file_inode (f=0xffff88000e61f700) at include/linux/fs.h:1963
1963     return f->f_inode;
>>>

```

read\_write.c security.c fsnotify.h »38

```

2378 static inline void file_start_write(struct file *file)
2379 {
2380     if (!S_ISREG(file_inode(file)->i_mode))
2381         return;
2382     _sb_start_write(file_inode(file)->i_sb, SB_FREEZE_WRITE, true);
2383 }
2384
2385 static inline bool file_start_write_trylock(struct file *file)
2386 {
2387     if (!S_ISREG(file_inode(file)->i_mode))
2388         return true;
2389     return _sb_start_write(file_inode(file)->i_sb, SB_FREEZE_WRITE, false);
2390 }
2391
2392 static inline void file_end_write(struct file *file)

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8128f229 in file\_start\_write+44 at include/linux/fs.h:2383

```

file_start_write (file=0xffff88000e61f700) at include/linux/fs.h:2383
2383     _sb_start_write(file_inode(file)->i_sb, SB_FREEZE_WRITE, true);
>>> 

```

read\_write.c fs.h security.c super.c »38

```

1958 extern void iput(struct inode *);
1959 extern int generic_update_time(struct inode *, struct timespec *, int
1960
1961 static inline struct inode *file_inode(const struct file *f)
1962 {
1963     return f->f_inode;
1964 }
1965
1966 /* /sys/fs */
1967 extern struct kobject *fs_kobj;
1968
1969 #define MAX_RW_COUNT (INT_MAX & PAGE_CACHE_MASK)
1970
1971 #define FLOCK_VERIFY_READ 1
1972 #define FLOCK_VERIFY_WRITE 2
1973

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8128f18d in file\_inode+12 at include/linux/fs.h:1963

```

file_inode (f=0xffff88000e61f700) at include/linux/fs.h:1963
1963     return f->f_inode;
>>> 

```

read\_write.c fs.h security.c super.c »39

```

1194 if (unlikely(sb->s_writers.frozen >= level)) {
1195     if (!wait)
1196         return 0;
1197     wait_event(sb->s_writers.wait_unfrozen,
1198                sb->s_writers.frozen < level);
1199 }
1200
1201 #ifdef CONFIG_LOCKDEP
1202     acquire_freeze_lock(sb, level, !wait, _RET_IP_);
1203 #endif
1204     percpu_counter_inc(&sb->s_writers.counter[level-1]);
1205     /*
1206     * Make sure counter is updated before we check for frozen.
1207     * freeze_super() first sets frozen and then checks the counter.
1208     */
1209     smp_mb();

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff812961c0 in \_\_sb\_start\_write+20 at fs/super.c:1194

```

__sb_start_write (sb=0xffff88000ec0d000, level=1, wait=true) at fs/super.c:1194
1194     if (unlikely(sb->s_writers.frozen >= level)) {
>>> 

```

read\_write.c fs.h security.c super.c »39

```

1194     if (unlikely(sb->s_writers.frozen >= level)) {
1195         if (!wait)
1196             return 0;
1197         wait_event(sb->s_writers.wait_unfrozen,
1198                     sb->s_writers.frozen < level);
1199     }
1200
1201 #ifdef CONFIG_LOCKDEP
1202     acquire_freeze_lock(sb, level, !wait, _RET_IP_);
1203 #endif
1204     percpu_counter_inc(&sb->s_writers.counter[level-1]);
1205     /*
1206      * Make sure counter is updated before we check for frozen.
1207      * freeze_super() first sets frozen and then checks the counter.
1208      */
1209     smp_mb();

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8129627c in \_\_sb\_start\_write+208 at fs/super.c:1204

```

1204     percpu_counter_inc(&sb->s_writers.counter[level-1]);
>>>

```

read\_write.c fs.h super.c percpu\_counter.h »40

```

123     fbc->count += amount;
124     preempt_enable();
125 }
126
127 static inline void
128 __percpu_counter_add(struct percpu_counter *fbc, s64 amount, s32 batch)
129 {
130     percpu_counter_add(fbc, amount);
131 }
132
133 static inline s64 percpu_counter_read(struct percpu_counter *fbc)
134 {
135     return fbc->count;
136 }
137
138 */

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81293bff in percpu\_counter\_inc+12 at include/linux/percpu

```

percpu_counter_inc (fbc=0xffff88000ec0d208) at include/linux/percpu_counter.h:166
166         percpu_counter_add(fbc, 1);
>>>

```

read\_write.c fs.h super.c percpu\_counter.h »40

```

43 s64 __percpu_counter_sum(struct percpu_counter *fbc);
44 int percpu_counter_compare(struct percpu_counter *fbc, s64 rhs);
45
46 static inline void percpu_counter_add(struct percpu_counter *fbc, s64 amount)
47 {
48     __percpu_counter_add(fbc, amount, percpu_counter_batch);
49 }
50
51 static inline s64 percpu_counter_sum_positive(struct percpu_counter *fbc)
52 {
53     s64 ret = __percpu_counter_sum(fbc);
54     return ret < 0 ? 0 : ret;
55 }
56
57 static inline s64 percpu_counter_sum(struct percpu_counter *fbc)
58 {

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81293bbe in percpu\_counter\_add+16 at include/linux/percpu

```

percpu_counter_add (fbc=0xffff88000ec0d208, amount=1) at include/linux/percpu_counter.h:48
48         __percpu_counter_add(fbc, amount, percpu_counter_batch);
>>>

```

File: percpu\_counter.c

```

80     count = __this_cpu_read(*fbc->counters) + amount;
81     if (count >= batch || count <= -batch) {
82         unsigned long flags;
83         raw_spin_lock_irqsave(&fbc->lock, flags);
84         fbc->count += count;
85         __this_cpu_sub(*fbc->counters, count - amount);
86         raw_spin_unlock_irqrestore(&fbc->lock, flags);
87     } else {
88         __this_cpu_add(*fbc->counters, amount);
89     }
90     preempt_enable();
91 }
92 EXPORT_SYMBOL(__percpu_counter_add);
93
94 /*
95 * Add up all the per-cpu counts, return the result. This is a more accurate

```

Console:

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff8158614d in __percpu_counter_add+19 at lib/percpu_counter.c:80
    percpu_counter_add (fbc=0xffff88000ec0d208, amount=1, batch=32) at lib/percpu_counter.c:80
80         count = __this_cpu_read(*fbc->counters) + amount;
>>>

```

File: percpu\_counter.h

```

300 static inline void __this_cpu_preempt_check(const char *op) { }
301 #endif
302
303 #define __pcpu_size_call_return(stem, variable)
304 ({ \
305     typeof(variable) pscr_ret__; \
306     __verify_pcpu_ptr(&(variable)); \
307     switch(sizeof(variable)) { \
308         case 1: pscr_ret__ = stem##1(variable); break; \
309         case 2: pscr_ret__ = stem##2(variable); break; \
310         case 4: pscr_ret__ = stem##4(variable); break; \
311         case 8: pscr_ret__ = stem##8(variable); break; \
312     default: \
313         __bad_size_call_parameter(); break; \
314     } \
315     pscr_ret__ . \

```

Console:

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff81585efd in __this_cpu_preempt_check+12 at include/linux/percpu.h:300
    __this_cpu_preempt_check (op=0xffffffff81ebafa0 "read") at include/linux/percpu-defs.h:300
300     static inline void __this_cpu_preempt_check(const char *op) { }
>>>

```

File: percpu-defs.h

```

80     count = __this_cpu_read(*fbc->counters) + amount;
81     if (count >= batch || count <= -batch) {
82         unsigned long flags;
83         raw_spin_lock_irqsave(&fbc->lock, flags);
84         fbc->count += count;
85         __this_cpu_sub(*fbc->counters, count - amount);
86         raw_spin_unlock_irqrestore(&fbc->lock, flags);
87     } else {
88         __this_cpu_add(*fbc->counters, amount);
89     }
90     preempt_enable();
91 }
92 EXPORT_SYMBOL(__percpu_counter_add);
93
94 /*
95 * Add up all the per-cpu counts, return the result. This is a more accurate

```

Console:

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff81586188 in __percpu_counter_add+78 at lib/percpu_counter.c:81
    percpu_counter_add (fbc=0xffff88000ec0d208, amount=1, batch=32) at lib/percpu_counter.c:81
81         if (count >= batch || count <= -batch) {
>>>

```

```
percpu_counter.c percpu_counter.h percpu-defs.h »43
80     count = __this_cpu_read(*fbc->counters) + amount;
81     if (count >= batch || count <= -batch) {
82         unsigned long flags;
83         raw_spin_lock_irqsave(&fbc->lock, flags);
84         fbc->count += count;
85         __this_cpu_sub(*fbc->counters, count - amount);
86         raw_spin_unlock_irqrestore(&fbc->lock, flags);
87     } else {
88         this_cpu_add(*fbc->counters, amount);
89     }
90     preempt_enable();
91 }
92 EXPORT_SYMBOL(__percpu_counter_add);
93 */
94 */
95 /* Add up all the per-cpu counts return the result. This is a more accurate
   * way to do it than summing up all the counters directly.
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff8158623e in \_\_percpu\_counter\_add+260 at lib/percpu\_counter.c:88

88 this\_cpu\_add(\*fbc->counters, amount);

>>>

```
super.c percpu_counter.c percpu_counter.h »43
1199 }
1200
1201 #ifdef CONFIG_LOCKDEP
1202     acquire_freeze_lock(sb, level, !wait, _RET_IP_);
1203 #endif
1204     percpu_counter_inc(&sb->s_writers.counter[level-1]);
1205     /*
1206     * Make sure counter is updated before we check for frozen.
1207     * freeze_super() first sets frozen and then checks the counter.
1208     */
1209     smp_mb();
1210     if (unlikely(sb->s_writers.frozen >= level)) {
1211         __sb_end_write(sb, level);
1212         goto retry;
1213     }
1214     return 1;
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff812962ad in \_\_sb\_start\_write+257 at fs/super.c:1209

\_sb\_start\_write (sb=0xffff88000ec0d000, level=1, wait=true) at fs/super.c:1209

1209 smp\_mb();

>>> █

```
super.c percpu_counter.c percpu_counter.h »43
1200
1201 #ifdef CONFIG_LOCKDEP
1202     acquire_freeze_lock(sb, level, !wait, _RET_IP_);
1203 #endif
1204     percpu_counter_inc(&sb->s_writers.counter[level-1]);
1205     /*
1206     * Make sure counter is updated before we check for frozen.
1207     * freeze_super() first sets frozen and then checks the counter.
1208     */
1209     smp_mb();
1210     if (unlikely(sb->s_writers.frozen >= level)) {
1211         __sb_end_write(sb, level);
1212         goto retry;
1213     }
1214     return 1;
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

}

[+] Threads

[1] id 1 from 0xffffffff812962b0 in \_\_sb\_start\_write+260 at fs/super.c:1210

1210 if (unlikely(sb->s\_writers.frozen >= level)) {

>>> █

```
super.c  percpcu_counter.  percpcu_counter. »43
1204     percpu_counter_inc(&sb->s_writers.counter[level-1]);
1205     /*
1206      * Make sure counter is updated before we check for frozen.
1207      * freeze_super() first sets frozen and then checks the counter.
1208      */
1209     smp_mb();
1210     if (unlikely(sb->s_writers.frozen >= level)) {
1211         __sb_end_write(sb, level);
1212         goto retry;
1213     }
1214     return 1;
1215 }
1216 EXPORT_SYMBOL(__sb_start_write);
1217
1218 /**
1219 * sh_wait_write - wait until all writers to given file system finish
1220 */

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff812962de in \_\_sb\_start\_write+306 at fs/super.c:1214

```
1214         return 1;
>>>
```

```
read_write.c  fs.h  super.c  percpcu_counter. »42
583     return -EINVAL;
584     if (unlikely(!access_ok(VERIFY_READ, buf, count)))
585         return -EFAULT;
586
587     ret = rw_verify_area(WRITE, file, pos, count);
588     if (ret >= 0) {
589         count = ret;
590         file_start_write(file);
591         if (file->f_op->write)
592             ret = file->f_op->write(file, buf, count, pos);
593         else if (file->f_op->ao_write)
594             ret = do_sync_write(file, buf, count, pos);
595         else
596             ret = new_sync_write(file, buf, count, pos);
597         if (ret > 0)
598             fsnotifyv_modify(file);

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81290723 in vfs\_write+191 at fs/read\_write.c:591

```
vfs_write (file=0xffff88000e61f700, buf=0x247fae0 "./syscall\nolder\hw03.syscall.a
591         if (file->f_op->write)
>>>
```

```
read_write.c  fs.h  super.c  percpcu_counter. »42
583     return -EINVAL;
584     if (unlikely(!access_ok(VERIFY_READ, buf, count)))
585         return -EFAULT;
586
587     ret = rw_verify_area(WRITE, file, pos, count);
588     if (ret >= 0) {
589         count = ret;
590         file_start_write(file);
591         if (file->f_op->write)
592             ret = file->f_op->write(file, buf, count, pos);
593         else if (file->f_op->ao_write)
594             ret = do_sync_write(file, buf, count, pos);
595         else
596             ret = new_sync_write(file, buf, count, pos);
597         if (ret > 0)
598             fsnotifyv_modify(file);

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg count = 10

[+]

Threads

[1] id 1 from 0xffffffff81290734 in vfs\_write+208 at fs/read\_write.c:592

```
592             ret = file->f_op->write(file, buf, count, pos);
>>>
```

```
521 EXPORT_SYMBOL(do_sync_write);
522
523 ssize_t new_sync_write(struct file *filp, const char __user *buf, size_t
524 {
525     struct iovec iov = { .iov_base = (void __user *)buf, .iov_len = len
526     struct kiocb kiocb;
527     struct iov_iter iter;
528     ssize_t ret;
529
530     init_sync_kiocb(&kiocb, filp);
531     kiocb.ki_pos = *ppos;
532     kiocb.ki_nbytes = len;
533     iov_iter_init(&iter, WRITE, &iov, 1, len);
534
535     ret = filp->f_op->write_iter(&kinch, &iter);
536 }
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8129045d in new\_sync\_write+39 at fs/read\_write.c:526

new\_sync\_write (filp=0xffff88000e61f700, buf=0x247fae0 "./syscall\nolder/hw03.sys"
526 struct iovec iov = { .iov\_base = (void \_\_user \*)buf, .iov\_len = len
>>>

```
521 EXPORT_SYMBOL(do_sync_write);
522
523 ssize_t new_sync_write(struct file *filp, const char __user *buf, size_t
524 {
525     struct iovec iov = { .iov_base = (void __user *)buf, .iov_len = len
526     struct kiocb kiocb;
527     struct iov_iter iter;
528     ssize_t ret;
529
530     init_sync_kiocb(&kiocb, filp);
531     kiocb.ki_pos = *ppos;
532     kiocb.ki_nbytes = len;
533     iov_iter_init(&iter, WRITE, &iov, 1, len);
534
535     ret = filp->f_op->write_iter(&kinch, &iter);
536 }
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg pos = 0xffff88000e5bff18: 37

[+]

Threads

[1] id 1 from 0xffffffff81290473 in new\_sync\_write+61 at fs/read\_write.c:531

init\_sync\_kiocb(&kiocb, filp);

>>>

```
45     size_t          ki_nbytes; /* copy of iocb->aio_nbytes */
46
47     struct list_head    ki_list;   /* the aio core uses this
48                           * for cancellation */
49
50     /*
51      * If the aio_resfd field of the userspace iocb is not zero,
52      * this is the underlying eventfd context to deliver events to.
53      */
54     struct eventfd_ctx *ki_eventfd;
55 };
56
57 static inline bool is_sync_kiocb(struct kiocb *kiocb)
58 {
59     return kiocb->ki_ctx == NULL;
60 }
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8128f128 in init\_sync\_kiocb+16 at include/linux/aio.h:67

init\_sync\_kiocb (kiocb=0xffff88000e5bfe28, filp=0xffff88000e61f700) at include/linux/aio.h:67
67 .ki\_obj.tsk = current,

>>>

read\_write.c

```

7 #ifndef __ASSEMBLY__
8 struct task_struct;
9
10 DECLARE_PER_CPU(struct task_struct *, current_task);
11
12 static __always_inline struct task_struct *get_current(void)
13 {
14     return this_cpu_read_stable(current_task);
15 }
16
17 #define current get_current()
18
19 #endif /* __ASSEMBLY__ */
20
21 #endif /* _ASM_X86_CURRENT_H */
22

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8128f128 in get\_current+0 at ./arch/x86/include/asm/curre

get\_current () at ./arch/x86/include/asm/current.h:14

14 return this\_cpu\_read\_stable(current\_task);

>>>

read\_write.c

```

62 static inline void init_sync_kiocb(struct kiocb *kiocb, struct file *filp)
63 {
64     *kiocb = (struct kiocb) {
65         .ki_ctx = NULL,
66         .ki_filp = filp,
67         .ki_obj.tsk = current,
68     };
69 }
70
71 /* prototypes */
72 #ifdef CONFIG_AIO
73 extern ssize_t wait_on_sync_kiocb(struct kiocb *iocb);
74 extern void aio_complete(struct kiocb *iocb, long res, long res2);
75 struct mm_struct;
76 extern void exit_aio(struct mm_struct *mm);
77 extern long do_in_submit(aio_context_t ctx_id, long nr

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8128f139 in init\_sync\_kiocb+33 at include/linux/aio.h:67

init\_sync\_kiocb (kiocb=0xffff88000e5bfe28, filp=0xffff88000e61f700) at include/linux/aio.h:67

.ki\_obj.tsk = current,

>>> █

read\_write.c

```

62 static inline void init_sync_kiocb(struct kiocb *kiocb, struct file *filp)
63 {
64     *kiocb = (struct kiocb) {
65         .ki_ctx = NULL,
66         .ki_filp = filp,
67         .ki_obj.tsk = current,
68     };
69 }
70
71 /* prototypes */
72 #ifdef CONFIG_AIO
73 extern ssize_t wait_on_sync_kiocb(struct kiocb *iocb);
74 extern void aio_complete(struct kiocb *iocb, long res, long res2);
75 struct mm_struct;
76 extern void exit_aio(struct mm_struct *mm);
77 extern long do_in_submit(aio_context_t ctx_id, long nr

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg ppos = 0xffff88000e5bff18: 37

[+]

Threads

[1] id 1 from 0xffffffff8128f13c in init\_sync\_kiocb+36 at include/linux/aio.h:64

64 \*kiocb = (struct kiocb) {

>>>

```

522 EXPORT_SYMBOL(do_sync_write);
523
524 ssize_t new_sync_write(struct file *filp, const char __user *buf, size_t
525 {
526     struct iovec iov = { .iov_base = (void __user *)buf, .iov_len = len };
527     struct kiocb kiocb;
528     struct iov_iter iter;
529     ssize_t ret;
530
531     init_sync_kiocb(&kiocb, filp);
532     kiocb.ki_pos = *ppos;
533     kiocb.ki_nbytes = len;
534     iov_iter_init(&iter, WRITE, &iov, 1, len);
535
536     ret = filp->f_op->write_iter(&kiocb, &iter);
537     if (-EIOCBQUEUED == ret)

```

Console [+] Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81290489 in new\_sync\_write+83 at fs/read\_write.c:532

new\_sync\_write (filp=0xffff88000e61f700, buf=0x247fae0 ./syscall\nolder/hw03.  
532        kiocb.ki\_pos = \*ppos;  
>>>

```

523
524 ssize_t new_sync_write(struct file *filp, const char __user *buf, size_t
525 {
526     struct iovec iov = { .iov_base = (void __user *)buf, .iov_len = len };
527     struct kiocb kiocb;
528     struct iov_iter iter;
529     ssize_t ret;
530
531     init_sync_kiocb(&kiocb, filp);
532     kiocb.ki_pos = *ppos;
533     kiocb.ki_nbytes = len;
534     iov_iter_init(&iter, WRITE, &iov, 1, len);
535
536     ret = filp->f_op->write_iter(&kiocb, &iter);
537     if (-EIOCBQUEUED == ret)
538         ret = wait_on_sync_kinch(&kinch);

```

Console [+] Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg pos = 0xffff88000e5bff18: 37

[+]

Threads

[1] id 1 from 0xffffffff81290497 in new\_sync\_write+97 at fs/read\_write.c:533

533        kiocb.ki\_nbytes = len;  
>>>

```

524 ssize_t new_sync_write(struct file *filp, const char __user *buf, size_t
525 {
526     struct iovec iov = { .iov_base = (void __user *)buf, .iov_len = len };
527     struct kiocb kiocb;
528     struct iov_iter iter;
529     ssize_t ret;
530
531     init_sync_kiocb(&kiocb, filp);
532     kiocb.ki_pos = *ppos;
533     kiocb.ki_nbytes = len;
534     iov_iter_init(&iter, WRITE, &iov, 1, len);
535
536     ret = filp->f_op->write_iter(&kiocb, &iter);
537     if (-EIOCBQUEUED == ret)
538         ret = wait_on_sync_kiocb(&kiocb);
539     *nmos = kinch.ki_mos;

```

Console [+] Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg pos = 0xffff88000e5bff18: 37

[+]

Threads

[1] id 1 from 0xffffffff812904a2 in new\_sync\_write+108 at fs/read\_write.c:534

534        iov\_iter\_init(&iter, WRITE, &iov, 1, len);  
>>> █

```
read_write.c [fs.h current.h aio.h iov_iter.c] »43
325     if (segment_eq(get_fs(), KERNEL_DS)) {
326         direction |= ITER_KVEC;
327         i->type = direction;
328         i->kvec = (struct kvec *)iov;
329     } else {
330         i->type = direction;
331         i->iov = iov;
332     }
333     i->nr_segs = nr_segs;
334     i->iov_offset = 0;
335     i->count = count;
336 }
337 EXPORT_SYMBOL(iov_iter_init);
338
339 static void memcpy_from_page(char *to, struct page *page, size_t offset, :
340 }
```

```
Console Registers Problems Executables Debugger Console Memory
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff81567b0e in iov_iter_init+27 at lib/iov_iter.c:325

iov_iter_init (i=0xfffff88000e5bfe00, direction=1, iov=0xfffff88000e5bfe80, nr_segs=1)
325     if (segment_eq(get_fs(), KERNEL_DS)) {
>>>
```

```
read_write.c [thread_info.h aio.h iov_iter.c] »44
162     ti = (void *)this_cpu_read_stable(kernel_stack) +
163             KERNEL_STACK_OFFSET - THREAD_SIZE;
164
165     return ti;
166 }
167
168 static inline unsigned long current_stack_pointer(void)
169 {
170     unsigned long sp;
171 #ifdef CONFIG_X86_64
172     asm("mov %%rsp,%0" : "=g" (sp));
173 #else
174     asm("mov %%esp,%0" : "=g" (sp));
175 #endif
176     return sp;
177 }
```

```
Console Registers Problems Executables Debugger Console Memory
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff81566887 in current_thread_info+8 at ./arch/x86/include/asm
current_thread_info () at ./arch/x86/include/asm/thread_info.h:162
162     ti = (void *)this_cpu_read_stable(kernel_stack) +
>>>
```

```
read_write.c [thread_info.h aio.h iov_iter.c] »44
325     if (segment_eq(get_fs(), KERNEL_DS)) {
326         direction |= ITER_KVEC;
327         i->type = direction;
328         i->kvec = (struct kvec *)iov;
329     } else {
330         i->type = direction;
331         i->iov = iov;
332     }
333     i->nr_segs = nr_segs;
334     i->iov_offset = 0;
335     i->count = count;
336 }
337 EXPORT_SYMBOL(iov_iter_init);
338
339 static void memcpy_from_page(char *to, struct page *page, size_t offset, :
340 }
```

```
Console Registers Problems Executables Debugger Console Memory
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff81567b43 in iov_iter_init+80 at lib/iov_iter.c:330

iov_iter_init (i=0xfffff88000e5bfe00, direction=1, iov=0xfffff88000e5bfe80, nr_segs=1)
330     i->type = direction;
>>>
```

read\_write.c

```

526     struct iovec iov = { .iov_base = (void __user *)buf, .iov_len = len
527     struct kiocb kiocb;
528     struct iov_iter iter;
529     ssize_t ret;
530
531     init_sync_kiocb(&kiocb, filp);
532     kiocb.ki_pos = *ppos;
533     kiocb.ki_nbytes = len;
534     iov_iter_init(&iter, WRITE, &iov, 1, len);
535
536     ret = filp->f_op->write_iter(&kiocb, &iter);
537     if (-EIOCBQUEUED == ret)
538         ret = wait_on_sync_kiocb(&kiocb);
539     *ppos = kiocb.ki_pos;
540     return ret;
541

```

Console

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff812904c9 in new_sync_write+147 at fs/read_write.c:536
new_sync_write (filp=0xffff88000e61f700, buf=0x247fae0 "./syscall\nolder/hw03.sysc"
536             ret = filp->f_op->write_iter(&kiocb, &iter);
>>>

```

read\_write.c

```

2649 *
2650 * This is a wrapper around __generic_file_write_iter() to be used by mc
2651 * filesystems. It takes care of syncing the file in case of O_SYNC file
2652 * and acquires i_mutex as needed.
2653 */
2654 ssize_t generic_file_write_iter(struct kiocb *iocb, struct iov_iter *from,
2655 {
2656     struct file *file = iocb->ki_filp;
2657     struct inode *inode = file->f_mapping->host;
2658     ssize_t ret;
2659
2660     mutex_lock(&inode->i_mutex);
2661     ret = __generic_file_write_iter(iocb, from);
2662     mutex_unlock(&inode->i_mutex);
2663
2664     if (ret > 0) {

```

Console

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff811f27bb in generic_file_write_iter+16 at mm/filemap.c:2656
generic_file_write_iter (iocb=0xffff88000e5bfe28, from=0xffff88000e5bfe00) at mm/file
2656         struct file *file = iocb->ki_filp;
>>>

```

read\_write.c

```

97     might_sleep();
98     /*
99      * The locking fastpath is the 1->0 transition from
100     * 'unlocked' into 'locked' state.
101    */
102    __mutex_fastpath_lock(&lock->count, __mutex_lock_slowpath);
103    mutex_set_owner(lock);
104 }

105 EXPORT_SYMBOL(mutex_lock);
106 #endif
107
108 static __always_inline void ww_mutex_lock_acquired(struct ww_mutex *ww,
109                                                 struct ww_acquire_ctx *ww_ctx)
110 {
111 #ifdef CONFIG_DEBUG_MUTEXES

```

Console

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff81d5a9dd in mutex_lock+12 at kernel/locking/mutex.c:97
mutex_lock (lock=0xffff88000e6a0d50) at kernel/locking/mutex.c:97
97         might_sleep();
>>>

```

```

read_write.c filemap.c mutex.c core.c >47
4206     if (should_resched()) {
4207         preempt_schedule_common();
4208         return 1;
4209     }
4210     return 0;
4211 }
4212 EXPORT_SYMBOL(_cond_resched);
4213 /*
4214 * __cond_resched_lock() - if a reschedule is pending, drop the given lock,
4215 * call schedule, and on return reacquire the lock.
4216 *
4217 * This works OK both with and without CONFIG_PREEMPT. We do strange low-
4218 * level operations here to prevent schedule() from being called twice (once via
4219 * spin_unlock(), once by hand).
4220 */

```

Console Registers Problems Executables Debugger Console Memory OS R

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81d59b81 in \_cond\_resched+8 at kernel/sched/core.c:4206

\_cond\_resched () at kernel/sched/core.c:4206

4206 if (should\_resched()) {

>>>

```

filemap.c mutex.c core.c preempt.h >48
90 */
91 * Returns true when we need to resched and can (barring IRQ state).
92 */
93 static __always_inline bool should_resched(void)
94 {
95     return unlikely(!raw_cpu_read_4(__preempt_count));
96 }
97
98 #ifdef CONFIG_PREEMPT
99     extern asmlinkage void __preempt_schedule(void);
100 # define __preempt_schedule() asm ("call __preempt_schedule")
101     extern asmlinkage void preempt_schedule(void);
102 # ifdef CONFIG_CONTEXT_TRACKING
103     extern asmlinkage void __preempt_schedule_context(void);
104 # define __preempt_schedule_context() asm ("call __preempt_schedule_c
105     extern asmlinkage void preempt_schedule_context(void).
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81d59b81 in should\_resched+0 at ./arch/x86/include/asm/pre

should\_resched () at ./arch/x86/include/asm/preempt.h:95

95 return unlikely(!raw\_cpu\_read\_4(\_\_preempt\_count));

>>>

```

filemap.c mutex.c core.c preempt.h >48
4206     if (should_resched()) {
4207         preempt_schedule_common();
4208         return 1;
4209     }
4210     return 0;
4211 }
4212 EXPORT_SYMBOL(_cond_resched);
4213 /*
4214 * __cond_resched_lock() - if a reschedule is pending, drop the given lock,
4215 * call schedule, and on return reacquire the lock.
4216 *
4217 * This works OK both with and without CONFIG_PREEMPT. We do strange low-
4218 * level operations here to prevent schedule() from being called twice (once via
4219 * spin_unlock(), once by hand).
4220 */

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81d59b9c in \_cond\_resched+35 at kernel/sched/core.c:4206

\_cond\_resched () at kernel/sched/core.c:4206

4206 if (should\_resched()) {

>>>

```

2840 }
2841
2842 static void __sched notrace preempt_schedule_common(void)
2843 {
2844     do {
2845         __preempt_count_add(PREEMPT_ACTIVE);
2846         __schedule();
2847         __preempt_count_sub(PREEMPT_ACTIVE);
2848
2849         /*
2850          * Check again in case we missed a preemption opportunity
2851          * between schedule and now.
2852         */
2853         barrier();
2854     } while (need_resched());
2855 }

```

Console [+] Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[1] id 1 from 0xffffffff81d599df in preempt\_schedule\_common+15 at kernel/sched/core.c:2845

preempt\_schedule\_common () at kernel/sched/core.c:2845

2845 \_\_preempt\_count\_add(PREEMPT\_ACTIVE);

>>>

```

67 * The various preempt_count add/sub methods
68 */
69
70 static __always_inline void __preempt_count_add(int val)
71 {
72     raw_cpu_add_4(__preempt_count, val);
73 }
74
75 static __always_inline void __preempt_count_sub(int val)
76 {
77     raw_cpu_add_4(__preempt_count, -val);
78 }
79
80 */
81 /* Because we keep PREEMPT_NEED_RESCHED set when we do _not_ need to reschedule
82 * a decrement which hits zero means we have no preempt count and should

```

Console [+] Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[1] id 1 from 0xffffffff81d599df in \_\_preempt\_count\_add+0 at ./arch/x86/include/asm/preempt.h:72

\_\_preempt\_count\_add (val=2097152) at ./arch/x86/include/asm/preempt.h:72

72 raw\_cpu\_add\_4(\_\_preempt\_count, val);

>>>

```

2840 }
2841
2842 static void __sched notrace preempt_schedule_common(void)
2843 {
2844     do {
2845         __preempt_count_add(PREEMPT_ACTIVE);
2846         __schedule();
2847         __preempt_count_sub(PREEMPT_ACTIVE);
2848
2849         /*
2850          * Check again in case we missed a preemption opportunity
2851          * between schedule and now.
2852         */
2853         barrier();
2854     } while (need_resched());
2855 }

```

Console [+] Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[1] id 1 from 0xffffffff81d59a0e in preempt\_schedule\_common+62 at kernel/sched/core.c:2846

preempt\_schedule\_common () at kernel/sched/core.c:2846

2846 \_\_schedule();

>>>

```

filemap.c [mutex.c] [core.c] [preempt.h] »48
2716     unsigned long *switch_count;
2717     struct rq *rq;
2718     int cpu;
2719
2720     preempt_disable();
2721     cpu = smp_processor_id();
2722     rq = cpu_rq(cpu);
2723     rCU_note_context_switch();
2724     prev = rq->curr;
2725
2726     schedule_debug(prev);
2727
2728     if (sched_feat(HRTICK))
2729         hrtick_clear(rq);
2730
2731 */

```

Console [Registers Problems Executables Debugger Console Memory]

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff811d59708 in \_schedule+8 at kernel/sched/core.c:2721

\_schedule () at kernel/sched/core.c:2721

2721 cpu = smp\_processor\_id();

>>>

```

filemap.c [mutex.c] [core.c] [preempt.h] [tree.c] »48
285     trace_rcu_utilization(TPS("Start context switch"));
286     rCU_sched_qs();
287     rCU_preempt_note_context_switch();
288     if (unlikely(raw_cpu_read(rCU_sched_qs_mask)))
289         rCU_momentary_dyntick_idle();
290     trace_rcu_utilization(TPS("End context switch"));
291 }
292 EXPORT_SYMBOL_GPL(rcu_note_context_switch);
293
294 */
295 * Register a quiescent state for all RCU flavors. If there is an
296 * emergency, invoke rCU_momentary_dyntick_idle() to do a heavy-weight
297 * dyntick-idle quiescent state visible to other CPUs (but only for those
298 * RCU flavors in desperate need of a quiescent state, which will normally
299 * be none of them). Either way, do a lightweight quiescent state for
300 * all RCU flavors

```

Console [Registers Problems Executables Debugger Console Memory]

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111d70c in rcu\_note\_context\_switch+8 at kernel/rcu/tree.c

rcu\_note\_context\_switch () at kernel/rcu/tree.c:285

285 trace\_rcu\_utilization(TPS("Start context switch"));

>>>

```

mutex.c [core.c] [preempt.h] [tree.c] [rculib.c] »49
20 TRACE_EVENT(rcu_utilization,
21     TP_PROTO(const char *s),
22     TP_ARGS(s),
23     TP_STRUCT_entry(
24         __field(const char *, s)
25     ),
26     TP_fast_assign(
27         __entry->s = s;
28     ),
29     TP_printk("%s", __entry->s)
30 );
31
32
33
34
35

```

Console [Registers Problems Executables Debugger Console Memory]

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111d37e in trace\_rcu\_utilization+20 at include/trace/events/trace\_rcu\_utilization (s=0xfffffff81e26aa2 "Start scheduler-tick") at include/trace/trace\_rcu\_utilization

20 TRACE\_EVENT(rcu\_utilization,

>>>

```

149 }
150
151 static __always_inline bool static_key_false(struct static_key *key)
152 {
153     if (unlikely(static_key_count(key) > 0))
154         return true;
155     return false;
156 }
157
158 static __always_inline bool static_key_true(struct static_key *key)
159 {
160     if (likely(static_key_count(key) > 0))
161         return true;
162     return false;
163 }
164

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111d37e in static\_key\_false+0 at include/linux/jump\_label.h:153

static\_key\_false (key=0xffffffff822c64c8 <\_tracepoint\_rcu\_utilization+8>) at include/linux/jump\_label.h:153  
    if (unlikely(static\_key\_count(key) > 0))

>>>

```

83
84 #include <linux/atomic.h>
85
86 static inline int static_key_count(struct static_key *key)
87 {
88     return atomic_read(&key->enabled);
89 }
90
91 #ifdef HAVE_JUMP_LABEL
92
93 #define JUMP_LABEL_TYPE_FALSE_BRANCH    0UL
94 #define JUMP_LABEL_TYPE_TRUE_BRANCH    1UL
95 #define JUMP_LABEL_TYPE_MASK          1UL
96
97 static
98 inline struct ignum_entry *ignum_label_get_entries(struct static_key *key)

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111d18c in static\_key\_count+12 at include/linux/jump\_label.h:88

static\_key\_count (key=0xffffffff822c64c8 <\_tracepoint\_rcu\_utilization+8>) at include/linux/jump\_label.h:88  
    return atomic\_read(&key->enabled);

>>>

```

27     return ACCESS_ONCE((v)->counter);
28 }
29
30 /**
31 * atomic_set - set atomic variable
32 * @v: pointer of type atomic_t
33 * @i: required value
34 *
35 * Atomically sets the value of @v to @i.
36 */
37 static inline void atomic_set(atomic_t *v, int i)
38 {
39     v->counter = i;
40 }
41
42 /**

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111cf4e in atomic\_read+12 at ./arch/x86/include/asm/atomic.h:27

atomic\_read (v=0xffffffff822c64c8 <\_tracepoint\_rcu\_utilization+8>) at ./arch/x86/include/asm/atomic.h:27  
    return ACCESS\_ONCE((v)->counter);

>>>

```

atomic.h core.c tree.c rcu.h jump_label.h »50
150
151 static __always_inline bool static_key_false(struct static_key *key)
152 {
153     if (unlikely(static_key_count(key) > 0))
154         return true;
155     return false;
156 }
157
158 static __always_inline bool static_key_true(struct static_key *key)
159 {
160     if (likely(static_key_count(key) > 0))
161         return true;
162     return false;
163 }
164
165 static inline void static_key_slow_inc(struct static_key *key)

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111d39e in static\_key\_false+32 at include/linux/jump\_lab

static\_key\_false (key=0xffffffff822c64c8 <\_tracepoint\_rcu\_utilization+8>) at inc

155 return false;

>>>

```

atomic.h core.c tree.c rcu.h jump_label.h »50
20 TRACE_EVENT(rcu_utilization,
21
22     TP_PROTO(const char *s),
23
24     TP_ARGS(s),
25
26     TP_STRUCT_entry(
27         __field(const char *, s)
28     ),
29
30     TP_fast_assign(
31         __entry->s = s;
32     ),
33
34     TP_printk("%s", __entry->s)
35 ).
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111d3a3 in trace\_rcu\_utilization+57 at include/trace/c

trace\_rcu\_utilization (s=0xffffffff81e26aa2 "Start scheduler-tick") at include,

20 TRACE\_EVENT(rcu\_utilization,

>>>

```

atomic.h core.c tree.c rcu.h jump_label.h »50
2522 * false, there is no point in invoking rCU_check_callbacks().
2523 */
2524 void rCU_check_callbacks(int user)
2525 {
2526     trace_rcu_utilization(TPS("Start scheduler-tick"));
2527     increment_cpu_stall_ticks();
2528     if (user || rCU_is_cpu_rrupt_from_idle()) {
2529
2530     /*
2531         * Get here if this CPU took its interrupt from user
2532         * mode or from the idle loop, and if this is not a
2533         * nested interrupt. In this case, the CPU is in
2534         * a quiescent state, so note it.
2535         *
2536         * No memory barrier is required here because both
2537         * rCU_sched_ns() and rCU_hh_ns() reference only CPU-local

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81121d94 in rCU\_check\_callbacks+26 at kernel/rcu/tree.c:25

rcu\_check\_callbacks (user=0) at kernel/rcu/tree.c:2527

2527 increment\_cpu\_stall\_ticks();

>>>

tree.c [+] rcu.h jump\_label.h tree\_plugin.h »52

```

1789     for_each_rcu_flavor(rsp)
1790         raw_cpu_inc(rsp->rda->ticks_this_gp);
1791     }
1792 
1793 #else /* #ifdef CONFIG_RCU_CPU_STALL_INFO */
1794 
1795 static void print_cpu_stall_info_begin(void)
1796 {
1797     pr_cont(" {");
1798 }
1799 
1800 static void print_cpu_stall_info(struct rcu_state *rsp, int cpu)
1801 {
1802     pr_cont(" %d", cpu);
1803 }
1804

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff811244e1 in increment\_cpu\_stall\_ticks+8 at kernel/rcu/tree.c:1789

increment\_cpu\_stall\_ticks () at kernel/rcu/tree\_plugin.h:1789

1789 for\_each\_rcu\_flavor(rsp)

>>>

tree.c [+] rcu.h jump\_label.h tree\_plugin.h »52

```

2522     * false, there is no point in invoking rCU_check_callbacks().
2523     */
2524 void rCU_check_callbacks(int user)
2525 {
2526     trace_rcu_utilization(TPS("Start scheduler-tick"));
2527     increment_cpu_stall_ticks();
2528     if (user || rCU_is_cpu_rrupt_from_idle()) {
2529 
2530     /*
2531         * Get here if this CPU took its interrupt from user
2532         * mode or from the idle loop, and if this is not a
2533         * nested interrupt. In this case, the CPU is in
2534         * a quiescent state, so note it.
2535         *
2536         * No memory barrier is required here because both
2537         * rCU_sched_qs() and rCU_hh_qs() reference only CPU-local

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81121d99 in rCU\_check\_callbacks+31 at kernel/rcu/tree.c:2528

rcu\_check\_callbacks (user=0) at kernel/rcu/tree.c:2528

2528 if (user || rCU\_is\_cpu\_rrupt\_from\_idle()) {

>>> █

tree.c [+] rcu.h jump\_label.h tree\_plugin.h »52

```

970     * interrupt from idle, return true. The caller must have at least
971     * disabled preemption.
972     */
973 static int rCU_is_cpu_rrupt_from_idle(void)
974 {
975     return __this_cpu_read(rcu_dynticks.dynticks_nesting) <= 1;
976 }
977 
978 /*
979  * Snapshot the specified CPU's dynticks counter so that we can later
980  * credit them with an implicit quiescent state. Return 1 if this CPU
981  * is in dynticks idle mode, which is an extended quiescent state.
982  */
983 static int dyntick_save_progress_counter(struct rcu_data *rdp,
984                                         bool *isidle, unsigned long *max)
985 {
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111e846 in rCU\_is\_cpu\_rrupt\_from\_idle+8 at kernel/rcu/tree.c:975

rcu\_is\_cpu\_rrupt\_from\_idle () at kernel/rcu/tree.c:975

975 return \_\_this\_cpu\_read(rcu\_dynticks.dynticks\_nesting) <= 1;

>>>

percpu-defs.h tree.c rcu.h tree\_plugin.h »s2

```

300 static inline void __this_cpu_preempt_check(const char *op) { }
301 #endif
302
303#define __pcpu_size_call_return(stem, variable)
304 ({ \
305     typeof(variable) pscr_ret_; \
306     _verify_pcpu_ptr(&(variable)); \
307     switch(sizeof(variable)) { \
308         case 1: pscr_ret_ = stem##1(variable); break; \
309         case 2: pscr_ret_ = stem##2(variable); break; \
310         case 4: pscr_ret_ = stem##4(variable); break; \
311         case 8: pscr_ret_ = stem##8(variable); break; \
312         default: \
313             __bad_size_call_parameter(); break; \
314     } \
315     pscr_ret_ . \

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111cc70 in \_\_this\_cpu\_preempt\_check+12 at include/linux/percpu-d

```

__this_cpu_preempt_check (op=0xffffffff81e26405 "read") at include/linux/percpu-d
300     static inline void __this_cpu_preempt_check(const char *op) { }
>>>

```

percpu-defs.h tree.c rcu.h tree\_plugin.h »s2

```

2540 */
2541
2542     rcu_sched_qs();
2543     rcu_bh_qs();
2544
2545 } else if (!in_softirq()) {
2546
2547/*
2548 * Get here if this CPU did not take its interrupt from
2549 * softirq, in other words, if it is not interrupting
2550 * a rCU_bh read-side critical section. This is an _bh
2551 * critical section, so note it.
2552 */
2553
2554     rcu_bh_qs();
2555

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81121db4 in rCU\_check\_callbacks+58 at kernel/rcu/tree.c:254

```

rcu_check_callbacks (user=0) at kernel/rcu/tree.c:254
2545     } else if (!in_softirq()) {
>>> s

```

percpu-defs.h preempt.h tree.c tree\_plugin.h »s2

```

17 * We mask the PREEMPT_NEED_RESCHED bit so as not to confuse all current
18 * that think a non-zero value indicates we cannot preempt.
19 */
20 static __always_inline int preempt_count(void)
21 {
22     return raw_cpu_read_4(__preempt_count) & ~PREEMPT_NEED_RESCHED;
23 }
24
25 static __always_inline void preempt_count_set(int pc)
26 {
27     raw_cpu_write_4(__preempt_count, pc);
28 }
29
30/*
31 * must be macros to avoid header recursion hell
32 */

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81121db4 in preempt\_count+0 at ./arch/x86/include/asm/pre

```

preempt_count () at ./arch/x86/include/asm/preempt.h:22
22     return raw_cpu_read_4(__preempt_count) & ~PREEMPT_NEED_RESCHED;
>>>

```

```

198     }
199 }
200
201 void rcu_bh_qs(void)
202 {
203     if (!__this_cpu_read(rcu_bh_data.passed_quiesce)) {
204         trace_rcu_grace_period(TPS("rcu_bh"),
205                                this_cpu_read(rcu_bh_data.gpnum),
206                                TPS("cpuqs"));
207         __this_cpu_write(rcu_bh_data.passed_quiesce, 1);
208     }
209 }
210
211 static DEFINE_PER_CPU(int, rcu_sched_qs_mask);
212
213 static DFTNPF PFR_CPII(struct rcu_dynticks rcu_dynticks) = {

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111d53e in rcu\_bh\_qs+8 at kernel/rcu/tree.c:203

```

rcu_bh_qs () at kernel/rcu/tree.c:203
203         if (!__this_cpu_read(rcu_bh_data.passed_quiesce)) {
>>>

```

```

300 static inline void __this_cpu_preempt_check(const char *op) { }
301 #endif
302
303 #define __percpu_size_call_return(stem, variable) \
304 ({ \
305     typeof(variable) pscr_ret; \
306     __verify_percpu_ptr(&(variable)); \
307     switch(sizeof(variable)) { \
308     case 1: pscr_ret = stem##1(variable); break; \
309     case 2: pscr_ret = stem##2(variable); break; \
310     case 4: pscr_ret = stem##4(variable); break; \
311     case 8: pscr_ret = stem##8(variable); break; \
312     default: \
313         __bad_size_call_parameter(); break; \
314     } \
315     pscr_ret . \

```

Console Registers Problems Executables Debugger Console Memory OS

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111cc70 in \_\_this\_cpu\_preempt\_check+12 at include/linux/percpu.h

```

__this_cpu_preempt_check (op=0xffffffff801e26405 "read") at include/linux/percpu-defs.h
300     static inline void __this_cpu_preempt_check(const char *op) { }
>>>

```

```

2551     * critical section, so note it.
2552     */
2553
2554     rcu_bh_qs();
2555 }
2556 rCU preempt_check_callbacks();
2557 if (rcu_pending())
2558     invoke_rcu_core();
2559 if (user)
2560     rCU_note_voluntary_context_switch(current);
2561 trace_rcu_utilization(TPS("End scheduler-tick"));
2562 }
2563
2564 /*
2565 * Scan the leaf rCU_node structures, processing dyntick state for any tl
2566 * have not yet encountered a quiescent state using the function specific

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81121dd7 in rCU\_check\_callbacks+93 at kernel/rcu/tree.c:255

```

rcu_check_callbacks (user=0) at kernel/rcu/tree.c:2556
2556         rCU preempt_check_callbacks();
>>>

```

percpu-defs.h preempt.h tree.c tree\_plugin.h »52

```

2551     * critical section, so note it.
2552     */
2553
2554     rcu_bh_qs();
2555 }
2556 rcu_preempt_check_callbacks();
2557 if (rcu_pending())
2558     invoke_rcu_core();
2559 if (user)
2560     rCU_note_voluntary_context_switch(current);
2561 trace_rcu_utilization(TPS("End_scheduler-tick"));
2562 }
2563 */
2564 /* Scan the leaf rCU_node structures, processing dyntick state for any tl
2565 * have not yet encountered a quiescent state using the function specif:

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81121ddc in rCU\_check\_callbacks+98 at kernel/rcu/tree.c:2557

rcu\_check\_callbacks (user=0) at kernel/rcu/tree.c:2557

2557 if (rcu\_pending())

>>> █

percpu-defs.h preempt.h tree.c tree\_plugin.h »52

```

3302 */
3303 static int rCU_pending(void)
3304 {
3305     struct rCU_state *rsp;
3306
3307     for_each_rcu_flavor(rsp)
3308         if (_rcu_pending(rsp, this_cpu_ptr(rsp->rda)))
3309             return 1;
3310     return 0;
3311 }
3312
3313 */
3314 * Return true if the specified CPU has any callback. If all_lazy is
3315 * non-NULL, store an indication of whether all callbacks are lazy.
3316 * (If there are no callbacks, all of them are deemed to be lazy.)
3317 */

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81123367 in rCU\_pending+8 at kernel/rcu/tree.c:3307

rcu\_pending () at kernel/rcu/tree.c:3307

3307 for\_each\_rcu\_flavor(rsp)

>>>

percpu-defs.h preempt.h tree.c tree\_plugin.h »52

```

3234 * carried out against CPU-local state are performed first. However,
3235 * we must check for CPU stalls first, else we might not get a chance.
3236 */
3237 static int __rcu_pending(struct rCU_state *rsp, struct rCU_data *rdp)
3238 {
3239     struct rCU_node *rnp = rdp->mynode;
3240
3241     rdp->n_rcu_pending++;
3242
3243     /* Check for CPU stalls, if enabled. */
3244     check_cpu_stall(rsp, rdp);
3245
3246     /* Is this CPU a NO_HZ_FULL CPU that should ignore RCU? */
3247     if (rcu_nohz_full_cpu(rsp))
3248         return 0;
3249

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff811230b8 in \_\_rcu\_pending+17 at kernel/rcu/tree.c:3239

\_rcu\_pending (rsp=0xffffffff8223d800 <rcu\_sched\_state>, rdp=0xffff88000fc144c0) at
3239 struct rCU\_node \*rnp = rdp->mynode;

>>> █

percpu-defs.h preempt.h tree.c tree\_plugin.h »52

```

3234 * carried out against CPU-local state are performed first. However,
3235 * we must check for CPU stalls first, else we might not get a chance.
3236 */
3237 static int __rcu_pending(struct rcu_state *rsp, struct rcu_data *rdp)
3238 {
3239     struct rcu_node *rnp = rdp->mynode;
3240
3241     rdp->n_rcu_pending++;
3242
3243     /* Check for CPU stalls, if enabled. */
3244     check_cpu_stall(rsp, rdp);
3245
3246     /* Is this CPU a NO_HZ_FULL CPU that should ignore RCU? */
3247     if (rcu_nohz_full_cpu(rsp))
3248         return 0;
3249

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

(no arguments)

[+]

Threads

[1] id 1 from 0xffffffff811230c4 in \_\_rcu\_pending+29 at kernel/rcu/tree.c:3241

```

3241         rdp->n_rcu_pending++;
>>>

```

percpu-defs.h preempt.h tree.c tree\_plugin.h »52

```

1261     unsigned long gps;
1262     unsigned long j;
1263     unsigned long js;
1264     struct rcu_node *rnp;
1265
1266     if (rcu_cpu_stall_suppress || !rcu_gp_in_progress(rsp))
1267         return;
1268     j = jiffies;
1269
1270     /*
1271      * Lots of memory barriers to reject false positives.
1272      *
1273      * The idea is to pick up rsp->gpnum, then rsp->jiffies_stall,
1274      * then rsp->gp_start, and finally rsp->completed. These values
1275      * are updated in the opposite order with memory barriers (or
1276      * equivalent) during grace-period initialization and cleanup.

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111f42f in check\_cpu\_stall+16 at kernel/rcu/tree.c:1266

```

check_cpu_stall (rsp=0xffffffff8223d800 <rcu_sched_state>, rdp=0xffff88000fc144c0)
1266         if (rcu_cpu_stall_suppress || !rcu_gp_in_progress(rsp))
>>>

```

percpu-defs.h preempt.h tree.c tree\_plugin.h »52

```

177     * permit this function to be invoked without holding the root rcu_node
178     * structure's ->lock, but of course results can be subject to change.
179 */
180 static int rcu_gp_in_progress(struct rcu_state *rsp)
181 {
182     return ACCESS_ONCE(rsp->completed) != ACCESS_ONCE(rsp->gpnum);
183 }
184
185 /*
186  * Note a quiescent state. Because we do not need to know
187  * how many quiescent states passed, just if there was at least
188  * one since the start of the grace period, this just sets a flag.
189  * The caller must have disabled preemption.
190 */
191 void rcu_sched_qs(void)
192 {

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111d4a7 in rcu\_gp\_in\_progress+12 at kernel/rcu/tree.c:182

```

rcu_gp_in_progress (rsp=0xffffffff8223d800 <rcu_sched_state>) at kernel/rcu/tree.c:182
182         return ACCESS_ONCE(rsp->completed) != ACCESS_ONCE(rsp->gpnum);
>>>

```

```

1262     unsigned long j;
1263     unsigned long js;
1264     struct rCU_node *rnp;
1265
1266     if (rcu_cpu_stall_suppress || !rcu_gp_in_progress(rsp))
1267         return;
1268     j = jiffies;
1269
1270     /*
1271      * Lots of memory barriers to reject false positives.
1272      *
1273      * The idea is to pick up rsp->gpnum, then rsp->jiffies stall,
1274      * then rsp->gp_start, and finally rsp->completed. These values
1275      * are updated in the opposite order with memory barriers (or
1276      * equivalent) during grace-period initialization and cleanup.
1277      * Now a false positive can occur if we get an new value of

```

Console [+] Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[1] id 1 from 0xffffffff8111f449 in check\_cpu\_stall+42 at kernel/rcu/tree.c:1267

check\_cpu\_stall (rsp=0xffffffff8223d800 <rcu\_sched\_state>, rdp=0xffff88000fc144c

1267 return;

>>>

```

3242     /* Check for CPU stalls, if enabled. */
3243     check_cpu_stall(rsp, rdp);
3244
3245     /* Is this CPU a NO_HZ FULL CPU that should ignore RCU? */
3246     if (rcu_nohz_full_cpu(rsp))
3247         return 0;
3248
3249     /* Is the RCU core waiting for a quiescent state from this CPU? */
3250     if (rcu_scheduler_fully_active &&
3251         rdp->qs_pending && !rdp->passed_quiesce &&
3252         rdp->rcu_qs_ctr_snap == __this_cpu_read(rcu_qs_ctr)) {
3253         rdp->n_rq_qs_pending++;
3254     } else if (rdp->qs_pending &&
3255                 (rdp->passed_quiesce ||
3256                  rdp->rcu_qs_ctr_snap != __this_cpu_read(rcu_qs_ctr))) {

```

Console [+] Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[1] id 1 from 0xffffffff811230f1 in \_\_rcu\_pending+74 at kernel/rcu/tree.c:324

\_\_rcu\_pending (rsp=0xffffffff8223d800 <rcu\_sched\_state>, rdp=0xffff88000fc144c

3247 if (rcu\_nohz\_full\_cpu(rsp))

>>>

```

3051     if (tick_nohz_full_cpu(smp_processor_id()) &&
3052         (!rcu_gp_in_progress(rsp) ||
3053          ULONG_CMP_LT(jiffies, ACCESS_ONCE(rsp->gp_start) + HZ)))
3054         return 1;
3055 #endif /* #ifdef CONFIG_NO_HZ_FULL */
3056     return 0;
3057 }
3058
3059 */
3060 * Bind the grace-period kthread for the sysidle flavor of RCU to the
3061 * timekeeping CPU.
3062 */
3063 static void rCU_bind_gp_kthread(void)
3064 {
3065     int __maybe_unused cpu;

```

Console [+] Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[1] id 1 from 0xffffffff81124725 in rCU\_nohz\_full\_cpu+12 at kernel/rcu/tree\_plu

rcu\_nohz\_full\_cpu (rsp=0xffffffff8223d800 <rcu\_sched\_state>) at kernel/rcu/tree

3056 return 0;

>>>

```

3242     /* Check for CPU stalls, if enabled. */
3243     check_cpu_stall(rsp, rdp);
3245
3246     /* Is this CPU a NO_HZ_FULL CPU that should ignore RCU? */
3247     if (rcu_nohz_full_cpu(rsp))
3248         return 0;
3249
3250     /* Is the RCU core waiting for a quiescent state from this CPU? */
3251     if (rcu_scheduler_fully_active &&
3252         rdp->qs_pending && !rdp->passed_quiesce &&
3253         rdp->rcu_qs_ctr_snap == __this_cpu_read(rcu_qs_ctr)) {
3254         rdp->n_rp_qs_pending++;
3255     } else if (rdp->qs_pending &&
3256                 (rdp->passed_quiesce ||
3257                  rdn->rcu_qs_ctr_snap != __this_cpu_read(rcu_qs_ctr))) {

```

Console [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8112310b in \_\_rcu\_pending+100 at kernel/rcu/tree.c:3251

```

_rcu_pending (rsp=0xffffffff8223d800 <rcu_sched_state>, rdp=0xffff88000fc144c0) at
3251     if (rcu_scheduler_fully_active &&
>>> 
```

```

495     * Does the CPU have callbacks ready to be invoked?
496     */
497 static int
498 cpu_has_callbacks_ready_to_invoke(struct rcu_data *rdp)
499 {
500     return &rdp->nxtlist != rdp->nxttail[RCU_DONE_TAIL] &&
501         rdp->nxttail[RCU_DONE_TAIL] != NULL;
502 }
503
504 /*
505  * Return the root node of the specified rcu_state structure.
506  */
507 static struct rcu_node *rcu_get_root(struct rcu_state *rsp)
508 {
509     return &rsp->node[0];
510 }

```

Console [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111d1d1 in cpu\_has\_callbacks\_ready\_to\_invoke+12 at kernel/

```

cpu_has_callbacks_ready_to_invoke (rdp=0xffff88000fc144c0) at kernel/rcu/tree.c:500
500     return &rdp->nxtlist != rdp->nxttail[RCU_DONE_TAIL] &&
>>> 
```

```

3264     rdp->n_rp_cb_ready++;
3265     return 1;
3266 }
3267
3268 /* Has RCU gone idle with this CPU needing another grace period? */
3269 if (cpu_needs_another_gp(rsp, rdp)) {
3270     rdp->n_rp_cpu_needs_gp++;
3271     return 1;
3272 }
3273
3274 /* Has another RCU grace period completed? */
3275 if (ACCESS_ONCE(rnp->completed) != rdp->completed) { /* outside lock */
3276     rdp->n_rp_gp_completed++;
3277     return 1;
3278 }
3279 
```

Console [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81123232 in \_\_rcu\_pending+395 at kernel/rcu/tree.c:3269

```

_rcu_pending (rsp=0xffffffff8223d800 <rcu_sched_state>, rdp=0xffff88000fc144c0) at
3269     if (cpu_needs_another_gp(rsp, rdp)) {
>>> 
```

File: tree.c

```

531 static int
532 cpu_needs_another_gp(struct rcu_state *rsp, struct rcu_data *rdp)
533 {
534     int i;
535
536     if (rcu_gp_in_progress(rsp))
537         return 0; /* No, a grace period is already in progress. */
538     if (rcu_future_needs_gp(rsp))
539         return 1; /* Yes, a no-CBs CPU needs one. */
540     if (!rdp->nxttail[RCU_NEXT_TAIL])
541         return 0; /* No, this is a no-CBs (or offline) CPU. */
542     if (*rdp->nxttail[RCU_NEXT_READY_TAIL])
543         return 1; /* Yes, this CPU has newly registered callbacks. */
544     for (i = RCU_WAIT_TAIL; i < RCU_NEXT_TAIL; i++)
545         if (rdp->nxttail[i - 1] != rdp->nxttail[i] &&
546             ULONG_CMP_LT(ACCESS_ONCE(rsp->completed),
547                         rdp->nxtcompleted[i]))
548             return 1;
549     return 0;
550 }

```

Console [C/C++ Attach to Application] gdb (8.1.0.20180409)

```

[+] Threads
[1] id 1 from 0xffffffff8111dad6 in cpu_needs_another_gp+16 at kernel/rcu/tree.c
cpu_needs_another_gp (rsp=0xffffffff8223d800 <rcu_sched_state>, rdp=0xffff88000f
536         if (rcu_gp_in_progress(rsp))
>>>

```

File: tree.c

```

177 * permit this function to be invoked without holding the root rcu_node
178 * structure's ->lock, but of course results can be subject to change.
179 */
180 static int rcu_gp_in_progress(struct rcu_state *rsp)
181 {
182     return ACCESS_ONCE(rsp->completed) != ACCESS_ONCE(rsp->gpnum);
183 }
184
185 /*
186 * Note a quiescent state. Because we do not need to know
187 * how many quiescent states passed, just if there was at least
188 * one since the start of the grace period, this just sets a flag.
189 * The caller must have disabled preemption.
190 */
191 void rcu_sched_qs(void)
192 {

```

Console [C/C++ Attach to Application] gdb (8.1.0.20180409)

```

[+] Threads
[1] id 1 from 0xffffffff8111d4a7 in rcu_gp_in_progress+12 at kernel/rcu/tree.c:1
rcu_gp_in_progress (rsp=0xffffffff8223d800 <rcu_sched_state>) at kernel/rcu/tree
182     return ACCESS_ONCE(rsp->completed) != ACCESS_ONCE(rsp->gpnum);
>>>

```

File: tree.c

```

533 {
534     int i;
535
536     if (rcu_gp_in_progress(rsp))
537         return 0; /* No, a grace period is already in progress. */
538     if (rcu_future_needs_gp(rsp))
539         return 1; /* Yes, a no-CBs CPU needs one. */
540     if (!rdp->nxttail[RCU_NEXT_TAIL])
541         return 0; /* No, this is a no-CBs (or offline) CPU. */
542     if (*rdp->nxttail[RCU_NEXT_READY_TAIL])
543         return 1; /* Yes, this CPU has newly registered callbacks. */
544     for (i = RCU_WAIT_TAIL; i < RCU_NEXT_TAIL; i++)
545         if (rdp->nxttail[i - 1] != rdp->nxttail[i] &&
546             ULONG_CMP_LT(ACCESS_ONCE(rsp->completed),
547                         rdp->nxtcompleted[i]))
548             return 1;
549     return 0;
550 }

```

Console [C/C++ Attach to Application] gdb (8.1.0.20180409)

```

[+] Threads
[1] id 1 from 0xffffffff8111daf0 in cpu_needs_another_gp+42 at kernel/rcu/tree.c:538
cpu_needs_another_gp (rsp=0xffffffff8223d800 <rcu_sched_state>, rdp=0xffff88000fc144
538         if (rcu_future_needs_gp(rsp))
>>>

```

```
514 * Interrupts must be disabled. If the caller does not hold the root
515 * rnp_node structure's ->lock, the results are advisory only.
516 */
517 static int rcu_future_needs_gp(struct rCU_state *rsp)
518 {
519     struct rCU_node *rnp = rCU_get_root(rsp);
520     int idx = (ACCESS_ONCE(rnp->completed) + 1) & 0x1;
521     int *fp = &rnp->need_future_gp[idx];
522
523     return ACCESS_ONCE(*fp);
524 }
525
526 /*
527 * Does the current CPU require a not-yet-started grace period?
528 * The caller must have disabled interrupts to prevent races with
529 * normal callback registry.
```

```
Console Registers Problems Executables Debugger Console Memory
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff8111da6b in rCU_future_needs_gp+12 at kernel/rCU/tree.c:5
rcu_future_needs_gp (rsp=0xffffffff8223d800 <rCU_sched_state>) at kernel/rCU/tree
519     struct rCU_node *rnp = rCU_get_root(rsp);
>>>
```

```
504 /*
505 * Return the root node of the specified rCU_state structure.
506 */
507 static struct rCU_node *rCU_get_root(struct rCU_state *rsp)
508 {
509     return &rsp->node[0];
510 }
511
512 /*
513 * Is there any need for future grace periods?
514 * Interrupts must be disabled. If the caller does not hold the root
515 * rnp_node structure's ->lock, the results are advisory only.
516 */
517 static int rCU_future_needs_gp(struct rCU_state *rsp)
518 {
519     struct rCU_node *rnp = rCU_get_root(rsp).
```

```
Console Registers Problems Executables Debugger Console Memory
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff8111da59 in rCU_get_root+12 at kernel/rCU/tree.c:509
rcu_get_root (rsp=0xffffffff8223d800 <rCU_sched_state>) at kernel/rCU/tree.c:509
509     return &rsp->node[0];
>>>
```

```
515 * rnp_node structure's ->lock, the results are advisory only.
516 */
517 static int rCU_future_needs_gp(struct rCU_state *rsp)
518 {
519     struct rCU_node *rnp = rCU_get_root(rsp);
520     int idx = (ACCESS_ONCE(rnp->completed) + 1) & 0x1;
521     int *fp = &rnp->need_future_gp[idx];
522
523     return ACCESS_ONCE(*fp);
524 }
525
526 /*
527 * Does the current CPU require a not-yet-started grace period?
528 * The caller must have disabled interrupts to prevent races with
529 * normal callback registry.
530 */
```

```
Console Registers Problems Executables Debugger Console Memory
debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff8111da7b in rCU_future_needs_gp+28 at kernel/rCU/tree.
rcu_future_needs_gp (rsp=0xffffffff8223d800 <rCU_sched_state>) at kernel/rCU/tree
520     int idx = (ACCESS_ONCE(rnp->completed) + 1) & 0x1;
>>>
```

```

535     if (rcu_gp_in_progress(rsp))
536         return 0; /* No, a grace period is already in progress. */
537     if (rcu_future_needs_gp(rsp))
538         return 1; /* Yes, a no-CBs CPU needs one. */
539     if (!rdp->nxttail[RCU_NEXT_TAIL])
540         return 0; /* No, this is a no-CBs (or offline) CPU. */
541     if (*rdp->nxttail[RCU_NEXT_READY_TAIL])
542         return 1; /* Yes, this CPU has newly registered callbacks. */
543     for (i = RCU_WAIT_TAIL; i < RCU_NEXT_TAIL; i++)
544         if (rdp->nxttail[i - 1] != rdp->nxttail[i] &&
545             ULONG_CMP_LT(ACCESS_ONCE(rsp->completed),
546                           rdp->nxtcompleted[i]))
547             return 1; /* Yes, CBs for future grace period. */
548     return 0; /* No grace period needed. */
549 }
550 }
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111db0a in cpu\_needs\_another\_gp+68 at kernel/rcu/tree.c:540

```

cpu_needs_another_gp (rsp=0xffffffff8223d800 <rcu_sched_state>, rdp=0xfffff8800000000)
540         if (!rdp->nxttail[RCU_NEXT_TAIL])
>>>
```

```

535     if (rcu_gp_in_progress(rsp))
536         return 0; /* No, a grace period is already in progress. */
537     if (rcu_future_needs_gp(rsp))
538         return 1; /* Yes, a no-CBs CPU needs one. */
539     if (!rdp->nxttail[RCU_NEXT_TAIL])
540         return 0; /* No, this is a no-CBs (or offline) CPU. */
541     if (*rdp->nxttail[RCU_NEXT_READY_TAIL])
542         return 1; /* Yes, this CPU has newly registered callbacks. */
543     for (i = RCU_WAIT_TAIL; i < RCU_NEXT_TAIL; i++)
544         if (rdp->nxttail[i - 1] != rdp->nxttail[i] &&
545             ULONG_CMP_LT(ACCESS_ONCE(rsp->completed),
546                           rdp->nxtcompleted[i]))
547             return 1; /* Yes, CBs for future grace period. */
548     return 0; /* No grace period needed. */
549 }
550 }
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111db38 in cpu\_needs\_another\_gp+114 at kernel/rcu/tree.c:544

```

544         for (i = RCU_WAIT_TAIL; i < RCU_NEXT_TAIL; i++)
>>> s
```

```

539     return 1; /* Yes, a no-CBs CPU needs one. */
540     if (!rdp->nxttail[RCU_NEXT_TAIL])
541         return 0; /* No, this is a no-CBs (or offline) CPU. */
542     if (*rdp->nxttail[RCU_NEXT_READY_TAIL])
543         return 1; /* Yes, this CPU has newly registered callbacks. */
544     for (i = RCU_WAIT_TAIL; i < RCU_NEXT_TAIL; i++)
545         if (rdp->nxttail[i - 1] != rdp->nxttail[i] &&
546             ULONG_CMP_LT(ACCESS_ONCE(rsp->completed),
547                           rdp->nxtcompleted[i]))
548             return 1; /* Yes, CBs for future grace period. */
549     return 0; /* No grace period needed. */
550 }

552 */
553 * rcu_eqs_enter_common - current CPU is moving towards extended quiesce
554 *
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111dbb0 in cpu\_needs\_another\_gp+234 at kernel/rcu/tree.c:549

```

549         return 0; /* No grace period needed. */
>>>
```

```

percpu-defs.h preempt.h tree.c tree_plugin.h "s2" (x)= N
3270     rdp->n_rp_cpu_needs_gp++;
3271     return 1;
3272 }
3273
3274 /* Has another RCU grace period completed? */
3275 if (ACCESS_ONCE(rnp->completed) != rdp->completed) { /* outside lock */
3276     rdp->n_rp_gp_completed++;
3277     return 1;
3278 }
3279
3280 /* Has a new RCU grace period started? */
3281 if (ACCESS_ONCE(rnp->gpnum) != rdp->gpnum ||
3282     unlikely(ACCESS_ONCE(rdp->gpwrap))) { /* outside lock */
3283     rdp->n_rp_gp_started++;
3284     return 1;
3285 }

```

Console Registers Problems Executables Debugger Console Memory OS Resources Perf Profiler

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8112326d in \_\_rcu\_pending+454 at kernel/rcu/tree.c:3275

\_rcu\_pending (rsp=0xffffffff8223d800 <rcu\_sched\_state>, rdp=0xffff88000fc144c0) at kernel/rcu/tree.c:3275

>>> s

```

percpu-defs.h preempt.h tree.c tree_plugin.h "s2" (x)= N
2581 {
2582 }
2583
2584 static int rcu_nocb_need_deferred_wakeup(struct rcu_data *rdp)
2585 {
2586     return false;
2587 }
2588
2589 static void do_nocb_deferred_wakeup(struct rcu_data *rdp)
2590 {
2591 }
2592
2593 static void rcu_spawn_all_nocb_kthreads(int cpu)
2594 {
2595 }
2596

```

Console Registers Problems Executables Debugger Console Memory OS Resources Perf Profiler

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81124671 in rcu\_nocb\_need\_deferred\_wakeup+12 at kernel/rcu/tree\_plugin.h:2586

rcu\_nocb\_need\_deferred\_wakeup (rdp=0xffff88000fc14640) at kernel/rcu/tree\_plugin.h:2586

>>> s

```

percpu-defs.h preempt.h tree.c tree_plugin.h "s2" (x)= Variables Br
3284     return 1;
3285 }
3286
3287 /* Does this CPU need a deferred NOCB wakeup? */
3288 if (rcu_nocb_need_deferred_wakeup(rdp)) {
3289     rdp->n_rp_nocb_defer_wakeup++;
3290     return 1;
3291 }
3292
3293 /* nothing to do */
3294 rdp->n_rp_need_nothing++;
3295 return 0;
3296 }
3297
3298 /* Check to see if there is any immediate RCU-related work to be done */
3299

```

Console Registers Problems Executables Debugger Console Memory OS Resources Perf Profiler

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81123339 in \_\_rcu\_pending+658 at kernel/rcu/tree.c:3294

\_rcu\_pending (rsp=0xffffffff8223de00 <rcu\_bh\_state>, rdp=0xffff88000fc14640) at kernel/rcu/tree.c:3294

>>> s

percpu-defs.h tree.c rcu.h tree\_plugin.h »s2

```

20 TRACE_EVENT(rcu_utilization,
21
22     TP_PROTO(const char *),
23
24     TP_ARGS(s),
25
26     TP_STRUCT_entry(
27         __field(const char *, s)
28     ),
29
30     TP_fast_assign(
31         __entry->s = s;
32     ),
33
34     TP_printk("%s", __entry->s)
35 ).
```

Console Registers Problems Executables Debugger Console Memory OS Resources Perf

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111d37e in trace\_rcu\_utilization+20 at include/trace/events/rcu.h:20

trace\_rcu\_utilization (s=0xffffffff81e26ab7 "End scheduler-tick") at include/trace/events/rcu.h:20

20 TRACE\_EVENT(rcu\_utilization,

>>> █

tree.c rcu.h jump\_label.h tree\_plugin.h »s2

```

150
151 static __always_inline bool static_key_false(struct static_key *key)
152 {
153     if (unlikely(static_key_count(key) > 0))
154         return true;
155     return false;
156 }
157
158 static __always_inline bool static_key_true(struct static_key *key)
159 {
160     if (likely(static_key_count(key) > 0))
161         return true;
162     return false;
163 }
164
165 static inline void static_key_slow_inc(struct static_key *key)
```

Console Registers Problems Executables Debugger Console Memory OS Resources Perf Profile V

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111d37e in static\_key\_false+0 at include/linux/jump\_label.h:153

static\_key\_false (key=0xffffffff822c64c8 <\_tracepoint\_rcu\_utilization+8>) at include/linux/jump\_label.h:153

153 if (unlikely(static\_key\_count(key) > 0))

>>> █

atomic.h tree.c rcu.h jump\_label.h »s2

```

27     return ACCESS_ONCE((v)->counter);
28 }
29
30 /**
31 * atomic_set - set atomic variable
32 * @v: pointer of type atomic_t
33 * @i: required value
34 *
35 * Atomically sets the value of @v to @i.
36 */
37 static inline void atomic_set(atomic_t *v, int i)
38 {
39     v->counter = i;
40 }
41
42 /**
```

Console Registers Problems Executables Debugger Console Memory OS Resources Perf Profi

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8111cf4e in atomic\_read+12 at ./arch/x86/include/asm/atomic.h:27

atomic\_read (v=0xffffffff822c64c8 <\_tracepoint\_rcu\_utilization+8>) at ./arch/x86/include/asm/atomic.h:27

27 return ACCESS\_ONCE((v)->counter);

>>> █

```

atomic.h
tree.c
rcu.h
jump_label.h
150
151 static __always_inline bool static_key_false(struct static_key *key)
152 {
153     if (unlikely(static_key_count(key) > 0))
154         return true;
155     return false;
156 }
157
158 static __always_inline bool static_key_true(struct static_key *key)
159 {
160     if (likely(static_key_count(key) > 0))
161         return true;
162     return false;
163 }
164
165 static inline void static_key_slow_inc(struct static_key *key)

```

Variables:

- key

Console:

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff8111d39e in static_key_false+32 at include/linux/jump_label.h:155
static_key_false (key=0xffffffff822c64c8 <_tracepoint_rcu_utilization+8>) at include/linux/jump_label.h:155
155     return false;
>>> 

```

```

atomic.h
tree.c
rcu.h
jump_label.h
20 TRACE_EVENT(rcu_utilization,
21
22     TP_PROTO(const char *s),
23
24     TP_ARGS(s),
25
26     TP_STRUCT_entry(
27         __field(const char *, s)
28     ),
29
30     TP_fast_assign(
31         __entry->s = s;
32     ),
33
34     TP_printk("%s", __entry->s)
35 );

```

Variables:

- s

Console:

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff8111d3a3 in trace_rcu_utilization+57 at include/trace/events/rcu.h:20
trace_rcu_utilization (s=0xffffffff81e26ab7 "End scheduler-tick") at include/trace/events/rcu.h:20
20     TRACE_EVENT(rcu_utilization,
>>> 

```

```

atomic.h
tree.c
rcu.h
jump_label.h
2554     rcu_bh_qs();
2555 }
2556     rcu_preempt_check_callbacks();
2557     if (rcu_pending())
2558         invoke_rcu_core();
2559     if (user)
2560         rcu_note_voluntary_context_switch(current);
2561     trace_rcu_utilization(TPS("End scheduler-tick"));
2562 }
2563
2564 /*
2565 * Scan the leaf rcu_node structures, processing dyntick state for any tl
2566 * have not yet encountered a quiescent state, using the function specif:
2567 * Also initiate boosting for any threads blocked on the root rcu_node.
2568 *
2569 * The caller must have suppressed start of new grace periods

```

Variables:

- s

Console:

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff811121e04 in rcu_check_callbacks+138 at kernel/rcu/tree.c:256
rcu_check_callbacks (user=0) at kernel/rcu/tree.c:2562
2562 }
>>> s 

```

atomic.h tree.c rcu.h jump\_label.h timer.c »52

```

1386     if (in_irq())
1387         irq_work_tick();
1388 #endif
1389     scheduler_tick();
1390     run_posix_cpu_timers(p);
1391 }
1392 */
1393 */
1394 /* This function runs timers and the timer-tq in bottom half context.
1395 */
1396 static void run_timer_softirq(struct softirq_action *h)
1397 {
1398     struct tvec_base *base = __this_cpu_read(tvec_bases);
1399
1400     hrtimer_run_pending();
1401 }
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8112baf1 in update\_process\_times+64 at kernel/time/timer.c:13

update\_process\_times (user tick=0) at kernel/time/timer.c:1386

1386 if (in\_irq())

>>>

preempt.h tree.c rcu.h timer.c »53

```

17 * We mask the PREEMPT_NEED_RESCHED bit so as not to confuse all current
18 * that think a non-zero value indicates we cannot preempt.
19 */
20 static __always_inline int preempt_count(void)
21 {
22     return raw_cpu_read_4(__preempt_count) & ~PREEMPT_NEED_RESCHED;
23 }
24
25 static __always_inline void preempt_count_set(int pc)
26 {
27     raw_cpu_write_4(__preempt_count, pc);
28 }
29
30 */
31 /* must be macros to avoid header recursion hell
32 */
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff8112baf1 in preempt\_count+0 at ./arch/x86/include/asm/preempt.h

preempt\_count () at ./arch/x86/include/asm/preempt.h:22

22 return raw\_cpu\_read\_4(\_\_preempt\_count) & ~PREEMPT\_NEED\_RESCHED;

>>>

preempt.h tree.c rcu.h timer.c irq\_work.c »53

```

173 }
174 EXPORT_SYMBOL_GPL(irq_work_run);
175
176 void irq_work_tick(void)
177 {
178     struct llist_head *raised = this_cpu_ptr(&raised_list);
179
180     if (!llist_empty(raised) && !arch_irq_work_has_interrupt())
181         irq_work_run_list(raised);
182     irq_work_run_list(this_cpu_ptr(&lazy_list));
183 }
184
185 */
186 /* Synchronize against the irq_work @entry, ensures the entry is not
187 * currently in use.
188 */
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff811d5784 in irq\_work\_tick+8 at kernel/irq\_work.c:178

irq\_work\_tick () at kernel/irq\_work.c:178

178 struct llist\_head \*raised = this\_cpu\_ptr(&raised\_list);

>>>

```
155 * test whether the list is empty without deleting something from the
156 * list.
157 */
158 static inline bool llist_empty(const struct llist_head *head)
159 {
160     return ACCESS_ONCE(head->first) == NULL;
161 }
162
163 static inline struct llist_node *llist_next(struct llist_node *node)
164 {
165     return node->next;
166 }
167
168 extern bool llist_add_batch(struct llist_node *new_first,
169                             struct llist_node *new_last,
170                             struct llist_head *head);
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff811d5073 in llist\_empty+12 at include/linux llist.h:160

```
llist_empty (head=0xfffff88000fc0e9a8) at include/linux(llist.h:160
160         return ACCESS_ONCE(head->first) == NULL;
>>>
```

```
1 timer.c 2 irq_work.c 3 llist.h 4 >5
67     return native_save_fl();
68 }
69
70 static inline notrace void arch_local_irq_restore(unsigned long flags)
71 {
72     native_restore_fl(flags);
73 }
74
75 static inline notrace void arch_local_irq_disable(void)
76 {
77     native_irq_disable();
78 }
79
80 static inline notrace void arch_local_irq_enable(void)
81 {
82     native_irq_enable();
83 }
```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff811d512c in arch\_local\_save\_flags+4 at ./arch/x86/include/as

```
arch_local_save_flags () at ./arch/x86/include/asm/irqflags.h:67
67         return native_save_fl();
>>> █
```

```
1 timer.c 2 irq_work.c 3 llist.h 4 >5
130 {
131     unsigned long flags;
132     struct irq_work *work;
133     struct llist_node *llnode;
134
135     BUG_ON(!irqs_disabled());
136
137     if (llist_empty(list))
138         return;
139
140     llnode = llist_del_all(list);
141     while (llnode != NULL) {
142         work = llist_entry(llnode, struct irq_work, llnode);
143
144         llnode = llist_next(llnode);
```

Console Registers Problems Executables Debugger Console Memory OS

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff811d5657 in irq\_work\_run\_list+48 at kernel/irq\_work.c:137

```
irq_work_run_list (list=0xfffff88000fc0e9b0) at kernel/irq_work.c:137
137         if (llist_empty(list))
>>>
```

irqFlags.h timer.c irq\_work.c llist.h "ss" (x)= Vari

```

155 * test whether the list is empty without deleting something from the
156 * list.
157 */
158 static inline bool llist_empty(const struct llist_head *head)
159 {
160     return ACCESS_ONCE(head->first) == NULL;
161 }
162
163 static inline struct llist_node *llist_next(struct llist_node *node)
164 {
165     return node->next;
166 }
167
168 extern bool llist_add_batch(struct llist_node *new_first,
169                             struct llist_node *new_last,
170                             struct llist_head *head);

```

Console Registers Problems Executables Debugger Console Memory OS Re

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff811d5073 in llist\_empty+12 at include/linux/list.h:160

llist\_empty (head=0xffff88000fc0e9b0) at include/linux/list.h:160

160 return ACCESS\_ONCE(head->first) == NULL;

>>>

irqFlags.h timer.c irq\_work.c llist.h "ss" (x)= Vari

```

130 {
131     unsigned long flags;
132     struct irq_work *work;
133     struct llist_node *llnode;
134
135     BUG_ON(!irqs_disabled());
136
137     if (llist_empty(list))
138         return;
139
140     llnode = llist_del_all(list);
141     while (llnode != NULL) {
142         work = llist_entry(llnode, struct irq_work, llnode);
143
144         llnode = llist_next(llnode);

```

Console Registers Problems Executables Debugger Console Memory OS Re

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff811d5667 in irq\_work\_run\_list+64 at kernel/irq\_work.c:138

irq\_work\_run\_list (list=0xffff88000fc0e9b0) at kernel/irq\_work.c:138

138 return;

>>>

irqFlags.h timer.c irq\_work.c llist.h "ss" (x)= Vari

```

1386     if (in_irq())
1387         irq_work_tick();
1388 #endif
1389     scheduler_tick();
1390     run_posix_cpu_timers(p);
1391 }
1392
1393 /*
1394 * This function runs timers and the timer-tq in bottom half context.
1395 */
1396 static void run_timer_softirq(struct softirq_action *h)
1397 {
1398     struct tvec_base *base = __this_cpu_read(tvec_bases);
1399
1400     hrtimer_run_pending();

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff8112bb14 in update\_process\_times+99 at kernel/time/time

update\_process\_times (user\_tick=0) at kernel/time/timer.c:1389

1389 scheduler\_tick();

>>>

irqflags.h core.c timer.c irq\_work.c llist.h »54

```

2465 * This function gets called by the timer code, with HZ frequency.
2466 * We call it with interrupts disabled.
2467 */
2468 void scheduler_tick(void)
2469 {
2470     int cpu = smp_processor_id();
2471     struct rq *rq = cpu_rq(cpu);
2472     struct task_struct *curr = rq->curr;
2473
2474     sched_clock_tick();
2475
2476     raw_spin_lock(&rq->lock);
2477     update_rq_clock(rq);
2478     curr->sched->task_tick(rq, curr, 0);
2479     update_cpu_load_active(rq);
2480     raw_spin_unlock(&rq->lock);

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff810d0721 in scheduler\_tick+8 at kernel/sched/core.c:2470

scheduler\_tick () at kernel/sched/core.c:2470

2470 int cpu = smp\_processor\_id();

>>>

core.c timer.c irq\_work.c llist.h clock.c »55

```

322 if (sched_clock_stable())
323     return;
324
325 if (unlikely(!sched_clock_running))
326     return;
327
328 WARN_ON_ONCE(!irqs_disabled());
329
330 scd = this_scd();
331 now_gtod = ktime_to_ns(ktime_get());
332 now = sched_clock();
333
334 scd->tick_raw = now;
335 scd->tick_gtod = now_gtod;
336 sched_clock_local(scd);
337

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff810d9c29 in sched\_clock\_tick+8 at kernel/sched/clock.c:322

sched\_clock\_tick () at kernel/sched/clock.c:322

322 if (sched\_clock\_stable())

>>>

core.c timer.c irq\_work.c llist.h clock.c »55

```

80 static struct static_key __sched_clock_stable = STATIC_KEY_INIT;
81 static int __sched_clock_stable_early;
82
83 int sched_clock_stable(void)
84 {
85     return static_key_false(&__sched_clock_stable);
86 }
87
88 static void __set_sched_clock_stable(void)
89 {
90     if (!sched_clock_stable())
91         static_key_slow_inc(&__sched_clock_stable);
92 }
93
94 void set_sched_clock_stable(void)
95 {

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff810d97a0 in sched\_clock\_stable+16 at kernel/sched/clock.c:85

sched\_clock\_stable () at kernel/sched/clock.c:85

85 return static\_key\_false(&\_\_sched\_clock\_stable);

>>>

core.c

```

150
151 static __always_inline bool static_key_false(struct static_key *key)
152 {
153     if (unlikely(static_key_count(key) > 0))
154         return true;
155     return false;
156 }
157
158 static __always_inline bool static_key_true(struct static_key *key)
159 {
160     if (likely(static_key_count(key) > 0))
161         return true;
162     return false;
163 }
164
165 static inline void static_key_slow_inc(struct static_key *key)

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff810d97a0 in static\_key\_false+0 at include/linux/jump\_label.h:153

```

static_key_false (key=0xffffffff825b4bd4 <__sched_clock_stable>) at include/linux/jump_label.h:153
153     if (unlikely(static_key_count(key) > 0))
>>>

```

atomic.h

```

27     return ACCESS_ONCE((v)->counter);
28 }
29
30 /**
31  * atomic_set - set atomic variable
32  * @v: pointer of type atomic_t
33  * @i: required value
34  *
35  * Atomically sets the value of @v to @i.
36  */
37 static inline void atomic_set(atomic_t *v, int i)
38 {
39     v->counter = i;
40 }
41
42 /**

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff810d9601 in atomic\_read+12 at ./arch/x86/include/asm/atomic

```

atomic_read (v=0xffffffff825b4bd4 <__sched_clock_stable>) at ./arch/x86/include/asm,
27         return ACCESS_ONCE((v)->counter);
>>>

```

atomic.h

```

62 #ifndef __ASSEMBLY__
63 #include <linux/types.h>
64
65 static inline notrace unsigned long arch_local_save_flags(void)
66 {
67     return native_save_fl();
68 }
69
70 static inline notrace void arch_local_irq_restore(unsigned long flags)
71 {
72     native_restore_fl(flags);
73 }
74
75 static inline notrace void arch_local_irq_disable(void)
76 {
77     native_irq_disable();

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff810d95ce in arch\_local\_save\_flags+4 at ./arch/x86/include/asm/irqflags.h:67

```

arch_local_save_flags () at ./arch/x86/include/asm/irqflags.h:67
67         return native_save_fl();
>>>

```

```

163     return raw_read_seqcount_begin(s);
164 }
165
166 /**
167 * raw_seqcount_begin - begin a seq-read critical section
168 * @s: pointer to seqcount_t
169 * Returns: count to be passed to read_seqcount_retry
170 *
171 * raw_seqcount_begin opens a read critical section of the given seqcount
172 * Validity of the critical section is tested by checking read_seqcount_r
173 * function.
174 *
175 * Unlike read_seqcount_begin(), this function will not wait for the cou
176 * to stabilize. If a writer is active when we begin, we will fail the
177 * read_seqcount_retry() instead of stabilizing at the beginning of the
178 * critical section

```

Console Registers Problems Executables Debugger Console Memory OS Reso

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff811356b1 in read\_seqcount\_begin+12 at include/linux/seqlock.h:16

```

read_seqcount_begin (s=0xffffffff825fac80 <tk_core>) at include/linux/seqlock.h:163
163     return raw_read_seqcount_begin(s);
>>>

```

```

141 * seqcount, but without any lockdep checking. Validity of the critical
142 * section is tested by checking read_seqcount_retry function.
143 */
144 static inline unsigned raw_read_seqcount_begin(const seqcount_t *s)
145 {
146     unsigned ret = __read_seqcount_begin(s);
147     smp_rmb();
148     return ret;
149 }
150
151 /**
152 * read_seqcount_begin - begin a seq-read critical section
153 * @s: pointer to seqcount_t
154 * Returns: count to be passed to read_seqcount_retry
155 *
156 * read_seqcount_begin opens a read critical section of the given seqcount

```

Console Registers Problems Executables Debugger Console Memory OS Reso

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff81135691 in raw\_read\_seqcount\_begin+12 at include/linux/seqlock.h:

```

raw_read_seqcount_begin (s=0xffffffff825fac80 <tk_core>) at include/linux/seqlock.h:146
146     unsigned ret = __read_seqcount_begin(s);
>>>

```

```

571 struct timekeeper *tk = &tk_core.timekeeper;
572 unsigned int seq;
573 ktime_t base;
574 s64 nsecs;
575
576 WARN_ON(timekeeping_suspended);
577
578 do {
579     seq = read_seqcount_begin(&tk_core.seq);
580     base = tk->tkr.base_mono;
581     nsecs = timekeeping_get_ns(&tk->tkr);
582
583 } while (read_seqcount_retry(&tk_core.seq, seq));
584
585
586 return ktime_add_ns(base, nsecs);

```

Console Registers Problems Executables Debugger Console Memory OS Reso

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff811367a4 in ktime\_get+84 at kernel/time/timekeeping.c:580

```

ktime_get () at kernel/time/timekeeping.c:580
580         base = tk->tkr.base_mono;
>>>

```

```

970     return (cycle_t)get_cycles();
971 }
972
973 /* .mask MUST be CLOCKSOURCE_MASK(64). See comment above read_tsc()
974 */
975 static struct clocksource clocksource_tsc = {
976     .name          = "tsc",
977     .rating        = 300,
978     .read          = read_tsc,
979     .mask          = CLOCKSOURCE_MASK(64),
980     .flags         = CLOCK_SOURCE_IS_CONTINUOUS |
981                      CLOCK_SOURCE_MUST_VERIFY,
982     .archdata      = { .vclock_mode = VCLOCK_TSC },
983 };
984
985

```

Console Registers Problems Executables Debugger Console Memory OS

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff81016e37 in read\_tsc+12 at arch/x86/kernel/tsc.c:970

read\_tsc (cs=0xffffffff82216a80 <clocksource\_tsc>) at arch/x86/kernel/tsc.c:970

970 return (cycle\_t)get\_cycles();

>>> S

```

117     asm volatile("rdtsc" : EAX_EDX_RET(val, low, high));
118
119     return EAX_EDX_VAL(val, low, high);
120 }
121
122 static inline unsigned long long native_read_pmc(int counter)
123 {
124     DECLARE_ARGS(val, low, high);
125
126     asm volatile("rdpmc" : EAX_EDX_RET(val, low, high) : "c" (counter));
127     return EAX_EDX_VAL(val, low, high);
128 }
129
130 #ifdef CONFIG_PARAVIRT
131 #include <asm/paravirt.h>
132 #else

```

Console Registers Problems Executables Debugger Console Memory OS

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff81015b20 in \_\_native\_read\_tsc+0 at ./arch/x86/include/asm/msr.h

\_\_native\_read\_tsc () at ./arch/x86/include/asm/msr.h:117

117 asm volatile("rdtsc" : EAX\_EDX\_RET(val, low, high));

>>> S

```

214     * retried).
215 */
216 static inline int read_seqcount_retry(const seqcount_t *s, unsigned start
217 {
218     smp_rmb();
219     return __read_seqcount_retry(s, start);
220 }
221
222
223
224 static inline void raw_write_seqcount_begin(seqcount_t *s)
225 {
226     s->sequence++;
227     smp_wmb();
228 }
229

```

Console Registers Problems Executables Debugger Console Memory OS Resources Perf

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff811356ee in read\_seqcount\_retry+15 at include/linux/seqlock.h:219

read\_seqcount\_retry (s=0xffffffff825fac80 <tk\_core>, start=34826) at include/linux/seqlock.h:219

219 return \_\_read\_seqcount\_retry(s, start);

>>> S

timekeeping.c

```

578     do {
579         seq = read_seqcount_begin(&tk_core.seq);
580         base = tk->tkr.base_mono;
581         nsecs = timekeeping_get_ns(&tk->tkr);
582
583     } while (read_seqcount_retry(&tk_core.seq, seq));
584
585     return ktime_add_ns(base, nsecs);
586 }
587 EXPORT_SYMBOL_GPL(ktime_get);
588
589 static ktime_t *offsets[TK_OFFSETS_MAX] = {
590     [TK_OFFSETS_REAL] = &tk_core.timekeeper.offsets_real,
591     [TK_OFFSETS_BOOT] = &tk_core.timekeeper.offsets_boot,
592     [TK_OFFSETS_TAI] = &tk_core.timekeeper.offsets_tai,
593 };

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff811367d5 in ktime\_get+133 at kernel/time/timekeeping.c:585

```

ktime_get () at kernel/time/timekeeping.c:585
585         return ktime_add_ns(base, nsecs);
>>>

```

tsc.c

```

136 #if defined(CONFIG_ARCH_SUPPORTS_INT128) && defined(__SIZEOF_INT128__)
137
138 #ifndef mul_u64_u32_shr
139 static inline u64 mul_u64_u32_shr(u64 a, u32 mul, unsigned int shift)
140 {
141     return (u64)((unsigned __int128)a * mul) >> shift;
142 }
143 #endif /* mul_u64_u32_shr */
144
145 #else
146
147 #ifndef mul_u64_u32_shr
148 static inline u64 mul_u64_u32_shr(u64 a, u32 mul, unsigned int shift)
149 {
150     u32 ah, al;
151     u64 ret;

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81015961 in mul\_u64\_u32\_shr+19 at include/linux/math64.h:141

```

mul_u64_u32_shr (a=134571950781, mul=465, shift=10) at include/linux/math64.h:141
141         return (u64)((unsigned __int128)a * mul) >> shift;
>>>

```

tsc.c

```

2466 * We call it with interrupts disabled.
2467 */
2468 void scheduler_tick(void)
2469 {
2470     int cpu = smp_processor_id();
2471     struct rq *rq = cpu_rq(cpu);
2472     struct task_struct *curr = rq->curr;
2473
2474     sched_clock_tick();
2475
2476     raw_spin_lock(&rq->lock);
2477     update_rq_clock(rq);
2478     curr->sched_class->task_tick(rq, curr, 0);
2479     update_cpu_load_active(rq);
2480     raw_spin_unlock(&rq->lock);
2481

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff810d0780 in scheduler\_tick+103 at kernel/sched/core.c:2476

```

scheduler_tick () at kernel/sched/core.c:2476
2476         raw_spin_lock(&rq->lock);
>>> s

```

spinlock.c

```

151     __raw_spin_lock(lock);
152 }
153 EXPORT_SYMBOL(__raw_spin_lock);
154 #endif
155
156 #ifndef CONFIG_INLINE_SPIN_LOCK_IRQSAVE
157 unsigned long __lockfunc __raw_spin_lock_irqsave(raw_spinlock_t *lock)
158 {
159     return __raw_spin_lock_irqsave(lock);
160 }
161 EXPORT_SYMBOL(__raw_spin_lock_irqsave);
162 #endif
163
164 #ifndef CONFIG_INLINE_SPIN_LOCK IRQ
165 void __lockfunc __raw_spin_lock_irq(raw_spinlock_t *lock)
166 {

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff81d5ce90 in \_\_raw\_spin\_lock+12 at kernel/locking/spinlock.c:1

```

__raw_spin_lock (lock=0xffff88000fc13a00) at kernel/locking/spinlock.c:151
151     __raw_spin_lock(lock);
>>> █

```

spinlock.c

```

145     LOCK_CONTENDED(lock, do_raw_spin_trylock, do_raw_spin_lock);
146 }
147
148 #endif /* !CONFIG_GENERIC_LOCKBREAK || CONFIG_DEBUG_LOCK_ALLOC */
149
150 static inline void __raw_spin_unlock(raw_spinlock_t *lock)
151 {
152     spin_release(&lock->dep_map, 1, _RET_IP_);
153     do_raw_spin_unlock(lock);
154     preempt_enable();
155 }
156
157 static inline void __raw_spin_unlock_irqrestore(raw_spinlock_t *lock,
158                                                 unsigned long flags)
159 {
160     spin_release(&lock->dep_map, 1, RFT_TP);

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff810f8f84e in \_\_raw\_spin\_lock+12 at include/linux/spinlock\_api.h

```

__raw_spin_lock (lock=0xffff88000fc13a00) at include/linux/spinlock_api.h:145
145     LOCK_CONTENDED(lock, do_raw_spin_trylock, do_raw_spin_lock);
>>> █

```

spinlock.c

```

675     return rq->cpu;
676 #else
677     return 0;
678 #endif
679 }
680
681 DECLARE_PER_CPU_SHARED_ALIGNED(struct rq, runqueues);
682
683 #define cpu_rq(cpu)      (&per_cpu(runqueues, (cpu)))
684 #define this_rq()         this_cpu_ptr(&runqueues)
685 #define task_rq(p)        cpu_rq(task_cpu(p))
686 #define cpu_curr(cpu)     (cpu_rq(cpu)->curr)
687 #define raw_rq()          raw_cpu_ptr(&runqueues)
688
689 static inline u64 __rq_clock_broken(struct rq *rq)
690 {

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff810c8d96 in cpu\_of+12 at kernel/sched/sched.h:675

```

cpu_of (rq=0xffff88000fc13a00) at kernel/sched/sched.h:675
675     return rq->cpu;
>>> █

```

File: tsc.c

```

7715 struct sched_entity *se = &curr->se;
7716
7717 for_each_sched_entity(se) {
7718     cfs_rq = cfs_rq_of(se);
7719     entity_tick(cfs_rq, se, queued);
7720 }
7721
7722 if (numbalancing_enabled)
7723     task_tick_numa(rq, curr);
7724
7725 update_rq_runnable_avg(rq, 1);
7726 }
7727
7728 /*
7729  * called on fork with the child task as argument from the parent's cont
7730  * - child not yet on the tasklist

```

Console:

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff810e51de in task_tick_fair+19 at kernel/sched/fair.c:7715
task_tick_fair (rq=0xffff88000fc13a00, curr=0xffff88000e630900, queued=0) at kernel/sched/fair.c:7715
7715     struct sched_entity *se = &curr->se;
>>> 

```

File: core.c

```

698 }
699
700 static inline u64 rq_clock_task(struct rq *rq)
701 {
702     lockdep_assert_held(&rq->lock);
703     return rq->clock_task;
704 }
705
706 #define RQCF_REQ_SKIP    0x01
707 #define RQCF_ACT_SKIP    0x02
708
709 static inline void rq_clock_skip_update(struct rq *rq, bool skip)
710 {
711     lockdep_assert_held(&rq->lock);
712     if (skip)
713         rq->clock_skip_update |= RQCF_RFO_SKTP;

```

Console:

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff810dc1d4 in rq_clock_task+12 at kernel/sched/sched.h:703
rq_clock_task (rq=0xffff88000fc13a00) at kernel/sched/sched.h:703
703     return rq->clock_task;
>>> 

```

File: sched.h

```

397 DEFINE_EVENT(sched_stat_runtime, sched_stat_runtime,
398     TP_PROTO(struct task_struct *tsk, u64 runtime, u64 vruntime),
399     TP_ARGS(tsk, runtime, vruntime));
400
401 /*
402  * Tracepoint for showing priority inheritance modifying a tasks
403  * priority.
404  */
405 TRACE_EVENT(sched_pi_setprio,
406
407     TP_PROTO(struct task_struct *tsk, int newprio),
408
409     TP_ARGS(tsk, newprio),
410
411     TP_STRUCT_entry(
412         array[ char ] comm
413         TASK_COMM_LEN )

```

Console:

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff810dc0d4 in trace_sched_stat_runtime+28 at include/trace/events/sched.h:397
trace_sched_stat_runtime (tsk=0xffff88000e630900, runtime=348195769, vruntime=18055250246) at include/trace/events/sched.h:397
397     DEFINE_EVENT(sched_stat_runtime, sched_stat_runtime,
>>> s

```

jump\_label.h    fair.c    sched.h    cpuacct.c    »61

```

227     { } /* terminate */
228 }
229
230/*
231 * charge this task's execution time to its accounting group.
232 *
233 * called with rq->lock held.
234 */
235 void cpuacct_charge(struct task_struct *tsk, u64 cputime)
236 {
237     struct cpuacct *ca;
238     int cpu;
239
240     cpu = task_cpu(tsk);
241
242     rCU_read_lock();

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff810f7103 in cpuacct\_charge+16 at kernel/sched/cpuacct.c:240

cpuacct\_charge (tsk=0xffff88000e630900, cputime=348195769) at kernel/sched/cpuacct.c:240  
240 cpu = task\_cpu(tsk);

>>>

fair.c    cpuacct.c    sched.h    rcupdate.h    »63

```

880     __rcu_read_lock();
881     _acquire(RCU);
882     rCU_lock_acquire(&rcu_lock_map);
883     rCU_lockdep_assert(rcu_is_watching(),
884                         "rcu_read_lock() used illegally while idle");
885 }
886
887/*
888 * So where is rCU_write_lock()? It does not exist, as there is no
889 * way for writers to lock out RCU readers. This is a feature, not
890 * a bug -- this property is what provides RCU's performance benefits.
891 * Of course, writers must coordinate with each other. The normal
892 * spinlock primitives work well for this, but any other technique may
893 * be used as well. RCU does not care how the writers keep out of each
894 * others' way, as long as they do so.
895 */

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff810f6984 in rCU\_read\_lock+4 at include/linux/rcupdate.h:880

rcu\_read\_lock () at include/linux/rcupdate.h:880  
880 \_\_rcu\_read\_lock();

>>>

Fair.c    cpuacct.c    sched.h    rcupdate.h    »63

```

234 */
235 void cpuacct_charge(struct task_struct *tsk, u64 cputime)
236 {
237     struct cpuacct *ca;
238     int cpu;
239
240     cpu = task_cpu(tsk);
241
242     rCU_read_lock();
243
244     ca = task_ca(tsk);
245
246     while (true) {
247         u64 *cpuusage = per_cpu_ptr(ca->cpuusage, cpu);
248         *cpuusage += cputime;
249

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff810f7117 in cpuacct\_charge+36 at kernel/sched/cpuacct.c:244

cpuacct\_charge (tsk=0xffff88000e630900, cputime=348195769) at kernel/sched/cpuacct.c:244  
244 ca = task\_ca(tsk);

>>>

```

cpuacct.c sched.h rcupdate.h cgroup.h »64
773     return task_css_check(task, subsys_id, false);
774 }
775
776 /**
777 * task_css_is_root - test whether a task belongs to the root css
778 * @task: the target task
779 * @subsys_id: the target subsystem ID
780 *
781 * Test whether @task belongs to the root css on the specified subsystem.
782 * May be invoked in any context.
783 */
784 static inline bool task_css_is_root(struct task_struct *task, int subsys_id)
785 {
786     return task_css_check(task, subsys_id, true) ==
787         init_css_set.subsys[subsys_id];
788 }

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff810f69dc in task\_css+15 at include/linux/cgroup.h:773

task\_css (task=0xffff88000e630900, subsys\_id=2) at include/linux/cgroup.h:773

773 return task\_css\_check(task, subsys\_id, false);

>>>

```

cpuacct.c sched.h rcupdate.h cgroup.h »64
240 }
241
242 static inline void __rcu_read_unlock(void)
243 {
244     preempt_enable();
245 }
246
247 static inline void synchronize_rcu(void)
248 {
249     synchronize_sched();
250 }
251
252 static inline int rcu_preempt_depth(void)
253 {
254     return 0;
255 }

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

**Threads**

[1] id 1 from 0xffffffff810f697e in \_\_rcu\_read\_unlock+4 at include/linux/rcupdate.h:245

\_\_rcu\_read\_unlock () at include/linux/rcupdate.h:245

245 }

>>> █

```

Fair.c cpuacct.c rcupdate.h stats.h »65
254 * running CPU and update the sum_exec_runtime field there.
255 */
256 static inline void account_group_exec_runtime(struct task_struct *tsk,
257                                               unsigned long long ns)
258 {
259     struct thread_group_cputimer *cputimer = &tsk->signal->cputimer;
260
261     if (!cputimer_running(tsk))
262         return;
263
264     raw_spin_lock(&cputimer->lock);
265     cputimer->cputime.sum_exec_runtime += ns;
266     raw_spin_unlock(&cputimer->lock);
267 }
268

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

}

[+]

**Threads**

[1] id 1 from 0xffffffff810dc424 in account\_group\_exec\_runtime+37 at kernel/sched/stats.h:261

261 if (!cputimer\_running(tsk))

>>> █

atomic-long.h sched.h

```

21     return ACCESS_ONCE((v)->counter);
22 }
23
24 /**
25  * atomic64_set - set atomic64 variable
26  * @v: pointer to type atomic64_t
27  * @i: required value
28 *
29 * Atomically sets the value of @v to @i.
30 */
31 static inline void atomic64_set(atomic64_t *v, long i)
32 {
33     v->counter = i;
34 }
35
36 /**

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff810db8cf in atomic64\_read+12 at ./arch/x86/include/asm/at

atomic64\_read (v=0xfffff88000fc13ae8) at ./arch/x86/include/asm/atomic64\_64.h:21

21 return ACCESS\_ONCE((v)->counter);

>>>

hrtimer.h atomic64\_64.h sched.h

```

397     return timer->state != HRTIMER_STATE_INACTIVE;
398 }
399
400 /**
401  * Helper function to check, whether the timer is on one of the queues
402  */
403 static inline int hrtimer_is_queued(struct hrtimer *timer)
404 {
405     return timer->state & HRTIMER_STATE_ENQUEUED;
406 }
407
408 /**
409  * Helper function to check, whether the timer is running the callback
410  * function
411  */
412 static inline int hrtimer_callback_running(struct hrtimer *timer)

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

arg queued = 0

[+] Threads

[1] id 1 from 0xffffffff810dbca8 in hrtimer\_active+29 at include/linux/hrtimer.h:3

398 }

>>> s

core.c sched.h thread\_info.h bitops.h

```

318     asm volatile("bt %2,%1\n\t"
319                 "sbb %0,%0"
320                 : "=r" (oldbit)
321                 : "m" (*(unsigned long *)addr), "Ir" (nr));
322
323     return oldbit;
324 }
325
326 #if 0 /* Fool kernel-doc since it doesn't do macros yet */
327 /**
328  * test_bit - Determine whether a bit is set
329  * @nr: bit number to test
330  * @addr: Address to start counting from
331  */
332 static int test_bit(int nr, const volatile unsigned long *addr);
333 #endif

```

Variables

Name
nr
addr
oldbit

Console Registers Problems Executables Debugger Console OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff810c6658 in variable\_test\_bit+16 at ./arch/x86/include/asm/bitops.h:318

variable\_test\_bit (nr=3, addr=0xfffff88000e5bc010) at ./arch/x86/include/asm/bitops.h:318

318 asm volatile("bt %2,%1\n\t"

>>> s

```

core.c sched.h thread_info.h bitops.h »67
91     return test_bit(flag, (unsigned long *)&ti->flags);
92 }
93
94 #define set_thread_flag(flag) \
95     set_ti_thread_flag(current_thread_info(), flag)
96 #define clear_thread_flag(flag) \
97     clear_ti_thread_flag(current_thread_info(), flag)
98 #define test_and_set_thread_flag(flag) \
99     test_and_set_ti_thread_flag(current_thread_info(), flag)
100 #define test_and_clear_thread_flag(flag) \
101     test_and_clear_ti_thread_flag(current_thread_info(), flag)
102 #define test_thread_flag(flag) \
103     test_ti_thread_flag(current_thread_info(), flag)
104
105 #define tif_need_resched() test_thread_flag(TIF_NEED_RESCHED)
106

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff810c7140 in test\_ti\_thread\_flag+39 at include/linux/thread\_info.h:92

test\_ti\_thread\_flag (ti=0xffff88000e5bc000, flag=3) at include/linux/thread\_info.h:92

92 }

>>>

```

core.c math64.h sched.h fair.c proc »67
466 static void __update_cpu_load(struct rq *this_rq, unsigned long this_load
467                                     unsigned long pending_updates)
468 {
469     int i, scale;
470
471     this_rq->nr_load_updates++;
472
473     /* Update our load: */
474     this_rq->cpu_load[0] = this_load; /* Fasttrack for idx 0 */
475     for (i = 1, scale = 2; i < CPU_LOAD_IDX_MAX; i++, scale += scale) {
476         unsigned long old_load, new_load;
477
478         /* scale is effectively 1 << i now, and >> i divides by scale */
479
480         old_load = this_rq->cpu_load[i];
481         old_load = decav_load_missed(old_load, pending_updates - 1, i);
482     }
483
484     this_rq->cpu_load[0] = this_load; /* Fasttrack for idx 0 */
485 }
486

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff810d9311 in \_\_update\_cpu\_load+40 at kernel/sched/proc.c:474

474 this\_rq->cpu\_load[0] = this\_load; /\* Fasttrack for idx 0 \*/

>>>

```

irqflags.h percpu-defs.h core.c core.c »69
300 static inline void __this_cpu preempt_check(const char *op) { }
301 #endif
302
303 #define __pcpu_size_call_return(stem, variable) \
304 ({ \
305     typeof(variable) pscr_ret_; \
306     __verify_pcpu_ptr(&(variable)); \
307     switch(sizeof(variable)) { \
308     case 1: pscr_ret_ = stem##1(variable); break; \
309     case 2: pscr_ret_ = stem##2(variable); break; \
310     case 4: pscr_ret_ = stem##4(variable); break; \
311     case 8: pscr_ret_ = stem##8(variable); break; \
312     default: \
313         __bad_size_call_parameter(); break; \
314     } \
315     pscr_ret_ . \

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+]

Threads

[1] id 1 from 0xffffffff811d873b in \_\_this\_cpu preempt\_check+12 at include/linux/percpu-defs.h:300

\_this\_cpu preempt\_check (op=0xffffffff81e39e81 "add") at include/linux/percpu-defs.h:300

300 static inline void \_\_this\_cpu preempt\_check(const char \*op) { }

>>>

profile.c x irq\_regs.h ptrace.h »75

```

411     struct pt_regs *regs = get_irq_regs();
412
413     if (!user_mode(regs) && prof_cpu_mask != NULL &&
414         cpumask_test_cpu(smp_processor_id(), prof_cpu_mask))
415         profile_hit(type, (void *)profile_pc(regs));
416 }
417
418 #ifdef CONFIG_PROC_FS
419 #include <linux/proc_fs.h>
420 #include <linux/seq_file.h>
421 #include <asm/uaccess.h>
422
423 static int prof_cpu_mask_proc_show(struct seq_file *m, void *v)
424 {
425     seq_printf(m, "%*pb\n", cpumask_pr_args(prof_cpu_mask));
426     return 0;

```

Console Registers Problems Executables Debugger Console Memory

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff81126401 in profile\_tick+36 at kernel/profile.c:414

profile\_tick (type=1) at kernel/profile.c:414  
414 cpumask\_test\_cpu(smp\_processor\_id(), prof\_cpu\_mask))  
>>>

bitops.h x profile.c ptrace.h cpumask.h »75

```

318     asm volatile("bt %2,%1\n\t"
319                 "sbb %0,%0"
320                 : "=r" (oldbit)
321                 : "m" (*(unsigned long *)addr), "Ir" (nr));
322
323     return oldbit;
324 }
325
326 #if 0 /* Fool kernel-doc since it doesn't do macros yet */
327 /**
328 * test_bit - Determine whether a bit is set
329 * @nr: bit number to test
330 * @addr: Address to start counting from
331 */
332 static int test_bit(int nr, const volatile unsigned long *addr);
333 #endif

```

Console Registers Problems Executables Debugger Console Memory OS Resources Perf Profile View

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff81124f3a in variable\_test\_bit+16 at ./arch/x86/include/asm/bitops.h:318

variable\_test\_bit (nr=0, addr=0xffffffff825f75b8 <prof\_cpu\_mask>) at ./arch/x86/include/asm/bitops.h:318  
318 asm volatile("bt %2,%1\n\t")  
>>>

hrtimer.c x tick-sched.c profile.c »76

```

807     u64 orun = 1;
808     ktime_t delta;
809
810     delta = ktime_sub(now, hrtimer_get_expires(timer));
811
812     if (delta.tv64 < 0)
813         return 0;
814
815     if (interval.tv64 < timer->base->resolution.tv64)
816         interval.tv64 = timer->base->resolution.tv64;
817
818     if (unlikely(delta.tv64 >= interval.tv64)) {
819         s64 incr = ktime_to_ns(interval);
820
821         orun = ktime_divns(delta, incr);
822         hrtimer_add_expires_ns(timer, incr * orun);

```

Console Registers Problems Executables Debugger Console Memory OS Resources

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)

[+] Threads

[1] id 1 from 0xffffffff8112dab9 in hrtimer\_forward+21 at kernel/time/hrtimer.c:807

hrtimer\_forward (timer=0xffff88000fc0d740, now=..., interval=...) at kernel/time/hrtimer.c:807  
807 u64 orun = 1;  
>>>

File: hrtimer.c

```

807     u64 orun = 1;
808     ktime_t delta;
809
810     delta = ktime_sub(now, hrtimer_get_expires(timer));
811
812     if (delta.tv64 < 0)
813         return 0;
814
815     if (interval.tv64 < timer->base->resolution.tv64)
816         interval.tv64 = timer->base->resolution.tv64;
817
818     if (unlikely(delta.tv64 >= interval.tv64)) {
819         s64 incr = ktime_to_ns(interval);
820
821         orun = ktime_divns(delta, incr);
822         hrtimer_add_expires_ns(timer, incr * orun);
823     }
824
825     return 0;
826 }

```

Console:

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff8112dac1 in hrtimer_forward+29 at kernel/time/hrtimer.c:810
810     delta = ktime_sub(now, hrtimer_get_expires(timer));
>>> s

```

File: read\_write.c

```

591     if (file->f_op->write)
592         ret = file->f_op->write(file, buf, count, pos);
593     else if (file->f_op->aio_write)
594         ret = do_sync_write(file, buf, count, pos);
595     else
596         ret = new_sync_write(file, buf, count, pos);
597     if (ret > 0) {
598         fsnotify_modify(file);
599         add_wchar(current, ret);
600     }
601     inc_syscw(current);
602     file_end_write(file);
603 }
604
605     return ret;
606 }

```

Console:

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff81290912 in SyS_write+20 at fs/read_write.c:635
SyS_write (fd=1, buf=17705776, count=35) at fs/read_write.c:635
635     SYSCALL_DEFINE3(write, unsigned int, fd, const char __user *, buf,
>>> s inc_syscw

```

File: read\_write.c

```

630     fdput_pos(f);
631 }
632
633     return ret;
634 }
635 SYSCALL_DEFINE3(write, unsigned int, fd, const char __user *, buf,
636                 size_t, count)
637 {
638     struct fd f = fdget_pos(fd);
639     ssize_t ret = -EBADF;
640
641     if (f.file) {
642         loff_t pos = file_pos_read(f.file);
643         ret = vfs_write(f.file, buf, count, &pos);
644         if (ret >= 0)
645             file_nos_write(f.file, ns);
646     }
647 }

```

Console:

```

debug kernel 4.0 [C/C++ Attach to Application] gdb (8.1.0.20180409)
[+]
Threads
[1] id 1 from 0xffffffff81290912 in SyS_write+20 at fs/read_write.c:635
SyS_write (fd=1, buf=17705776, count=35) at fs/read_write.c:635
635     SYSCALL_DEFINE3(write, unsigned int, fd, const char __user *, buf,
>>> s file_end_write

```

The screenshot shows a debugger interface with several windows:

- Code Editor:** Shows the file `fs/read_write.c` with lines 589 to 604. The code handles file writes, including sync and aio operations.
- Registers:** Tab is visible in the toolbar.
- Problems:** Tab is visible in the toolbar.
- Executables:** Tab is visible in the toolbar.
- Debugger Console:** Active tab, showing a stack trace for a kernel thread (id 1) at address `0xffffffff81290945` in `SYSC_write+19` at `fs/read_write.c:638`. The trace also includes the instruction `SYSC write (fd=1, buf=0x10e2b30 "使用 'syscall' 呼叫system call\n\n", count=35`.
- Mappings:** Tab is visible in the toolbar.